

Fundamentals of Image Processing

Review – Spatial Filtering

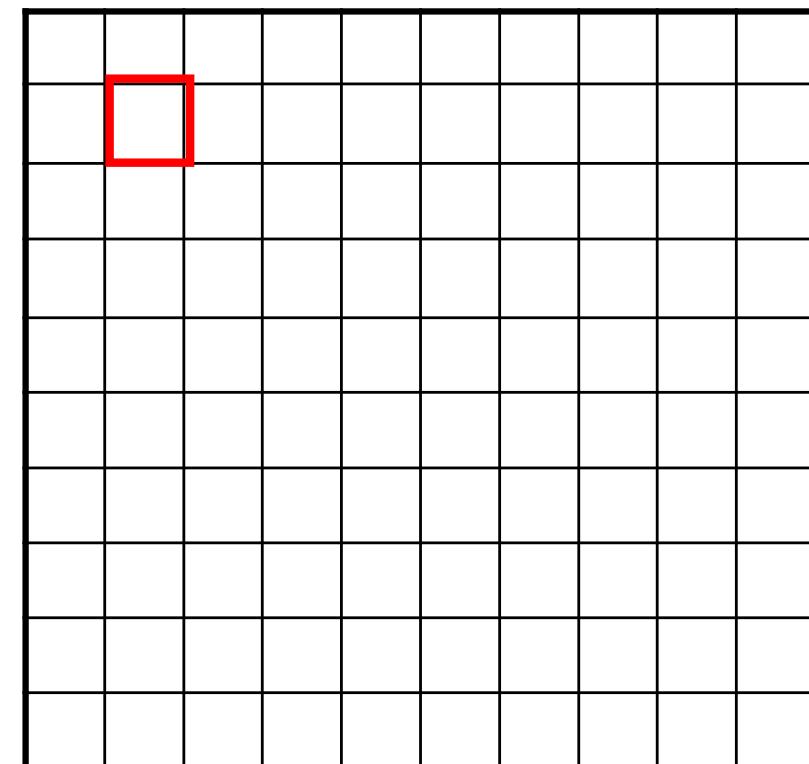
$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$



$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Review – Spatial Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

0	10									

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Review – Spatial Filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot]$$

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Review – Spatial Filtering

$$g[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Review – Spatial Filtering

$$g[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30					

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Review – Spatial Filtering

$$g[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Slide credit: S. Seitz

Why does a lower resolution image still make sense to us? What do we lose?

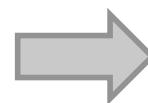


Image: <http://www.flickr.com/photos/igorms/136916757/>

Slide credit: D. Hoiem

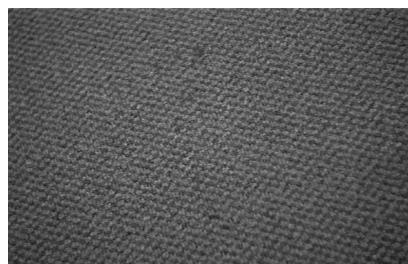
How is it that a 4MP image can be compressed to a few hundred KB without a noticeable change?



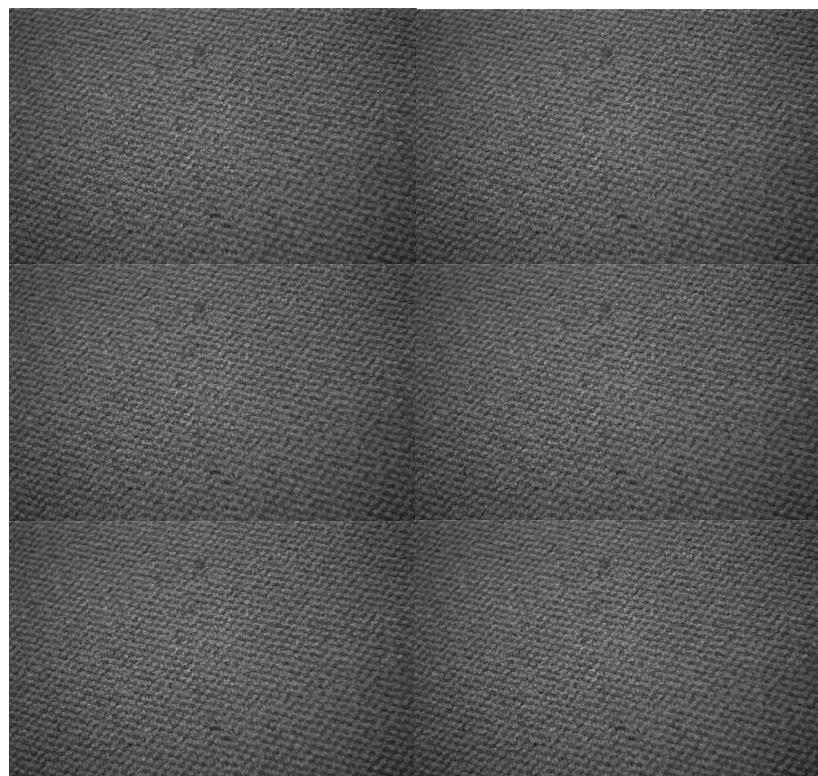
Slide credit: J. Hays

Answer to these questions?

- Thinking images in terms of frequency.
- Treat images as infinite-size, continuous periodic functions.



= ...



...

...

...

Jean Baptiste Joseph Fourier (1768-1830)

had idea (1807):

Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.



Slide credit: A. Efros

Jean Baptiste Joseph Fourier (1768-1830)

had an idea (1807):

Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.

- Don't believe it?
 - Neither did Lagrange, Laplace, Poisson and other big wigs
 - Not translated into English until 1878!
- But it's (mostly) true!
 - called Fourier Series
 - there are some subtle restrictions



Slide credit: A. Efros

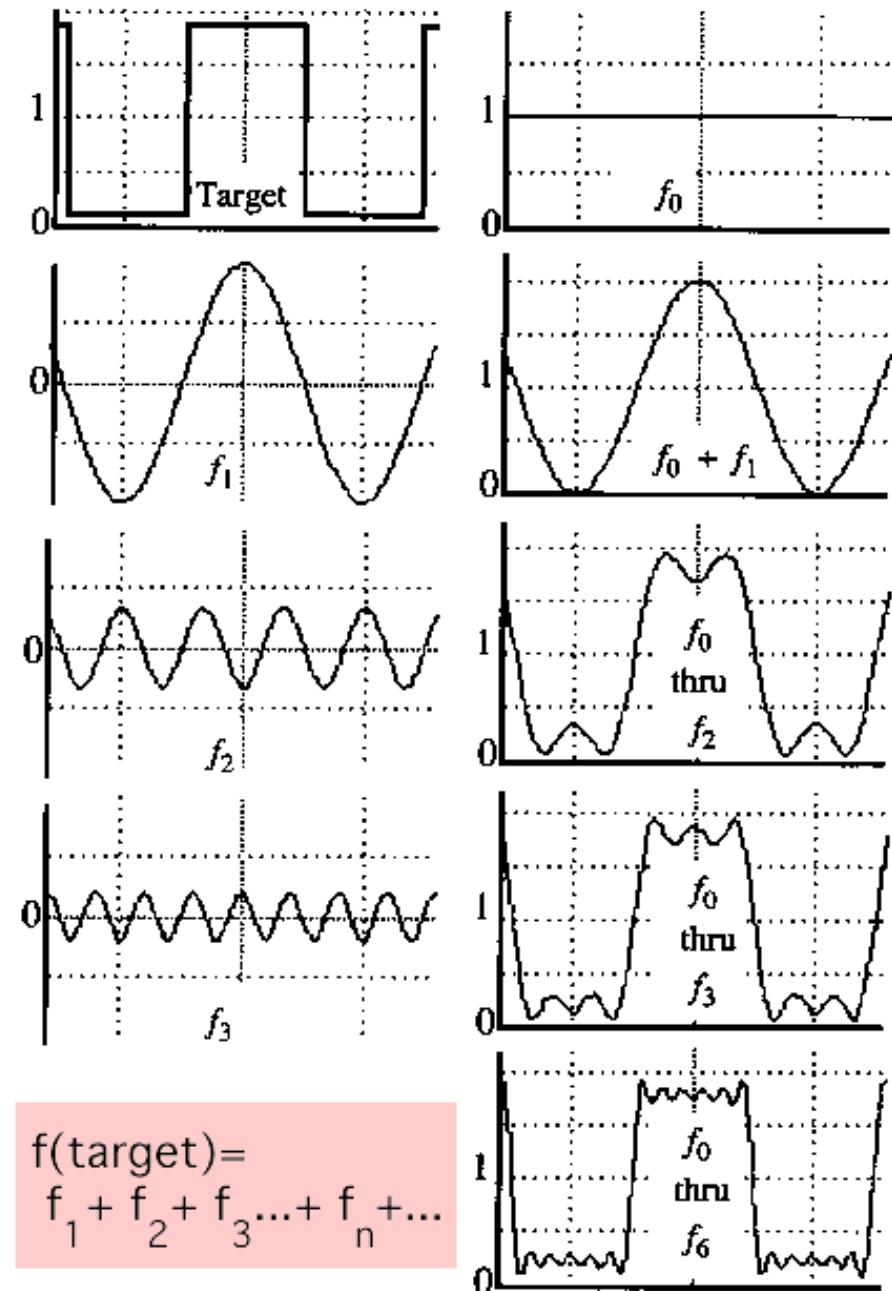
A sum of sines

Our building block:

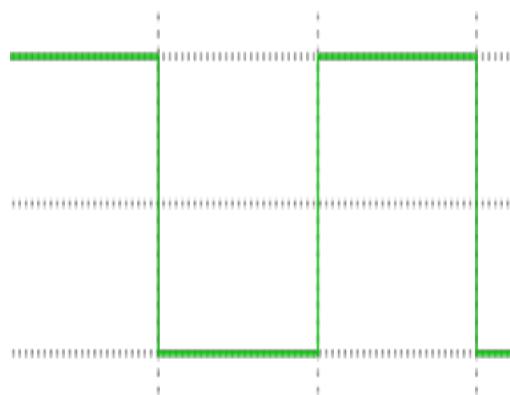
$$A \sin(\omega x + \phi)$$

The Fourier theory shows how most real function can be represented in terms of a basis of sinusoids.

Add enough of them to get any signal $f(x)$ you want.

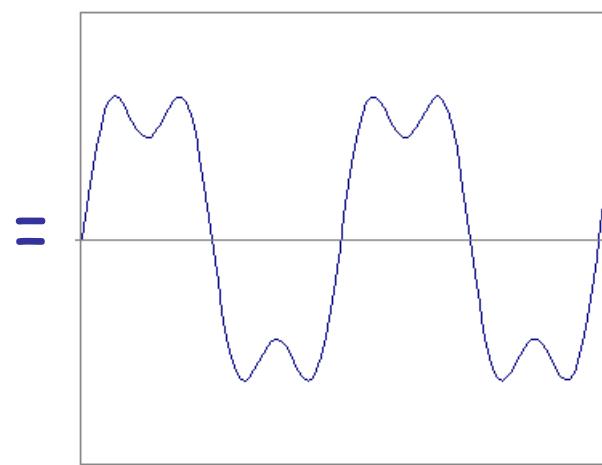
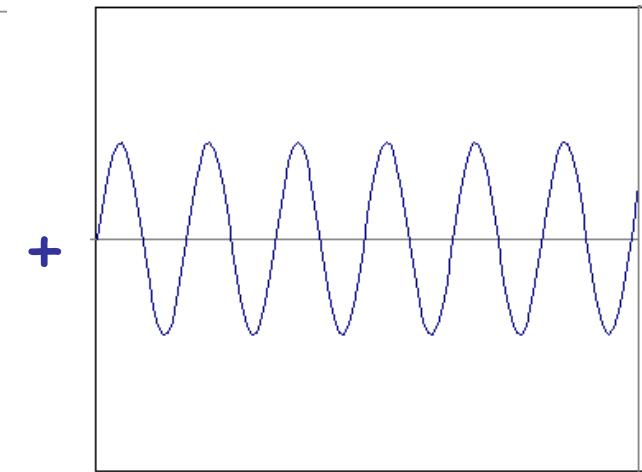
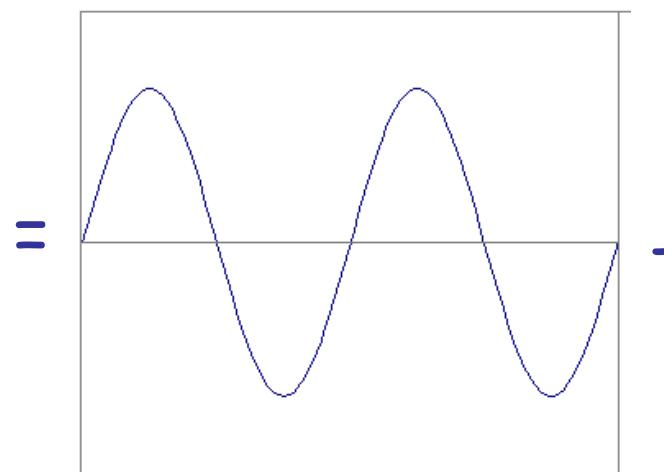
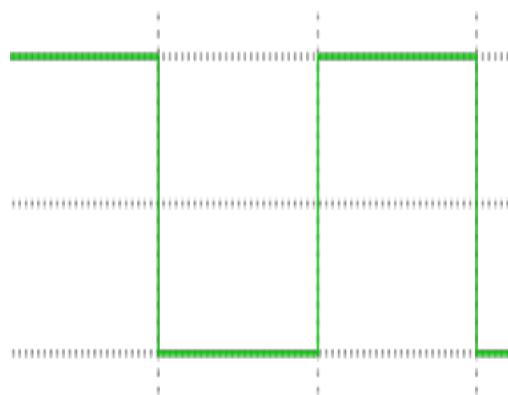


Frequency Spectra



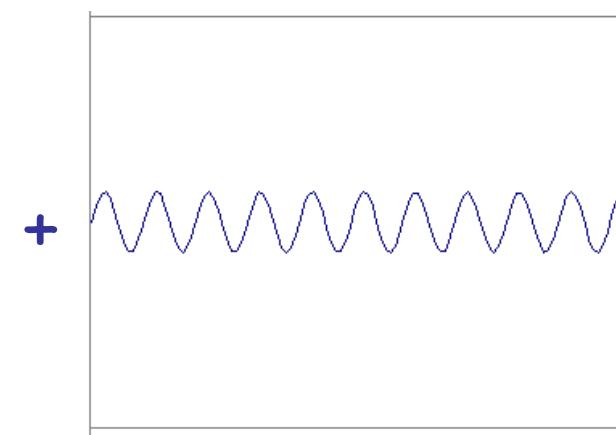
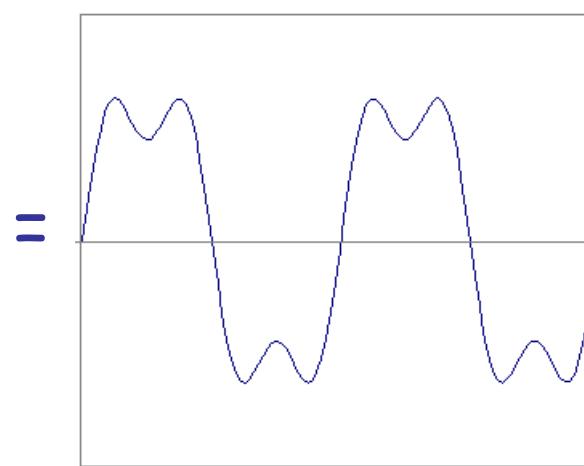
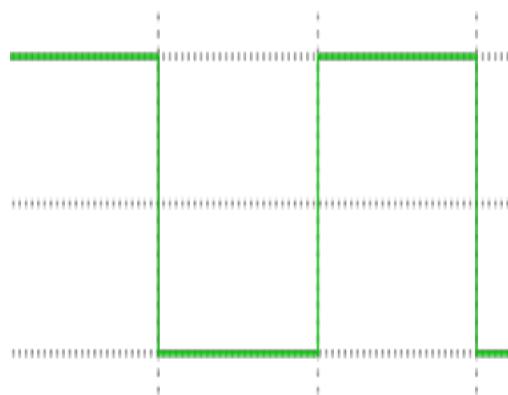
Slide credit: A. Efros

Frequency Spectra



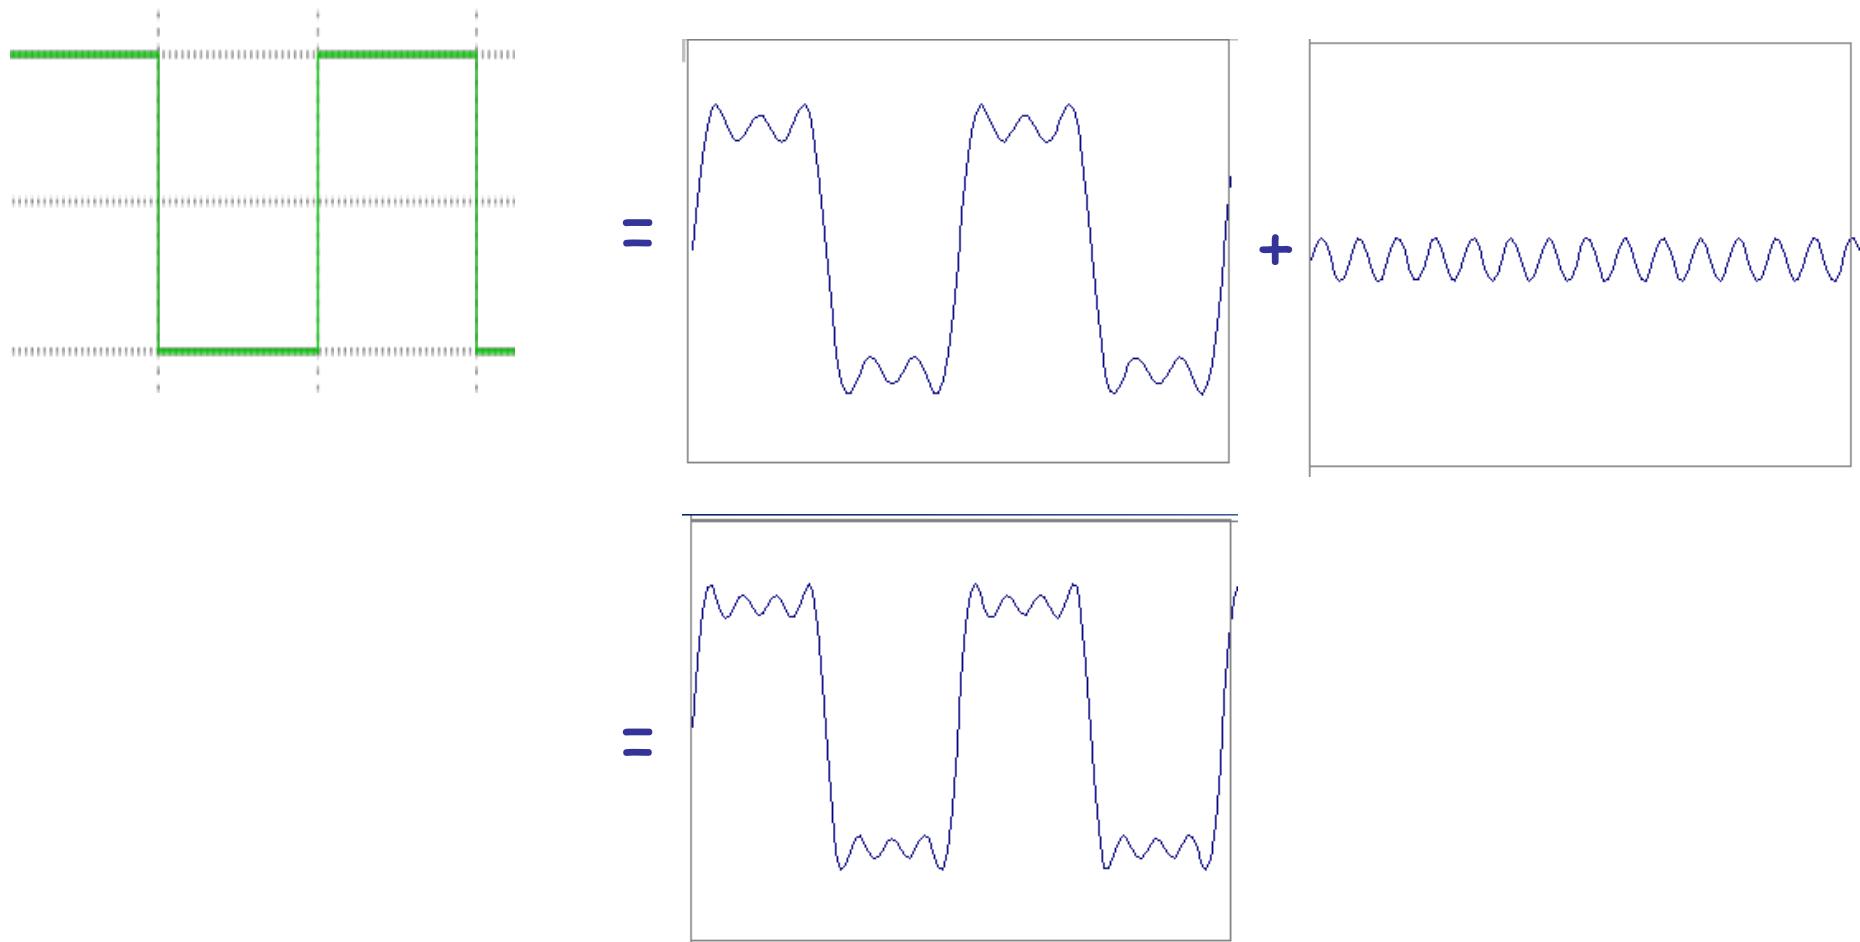
Slide credit: A. Efros

Frequency Spectra



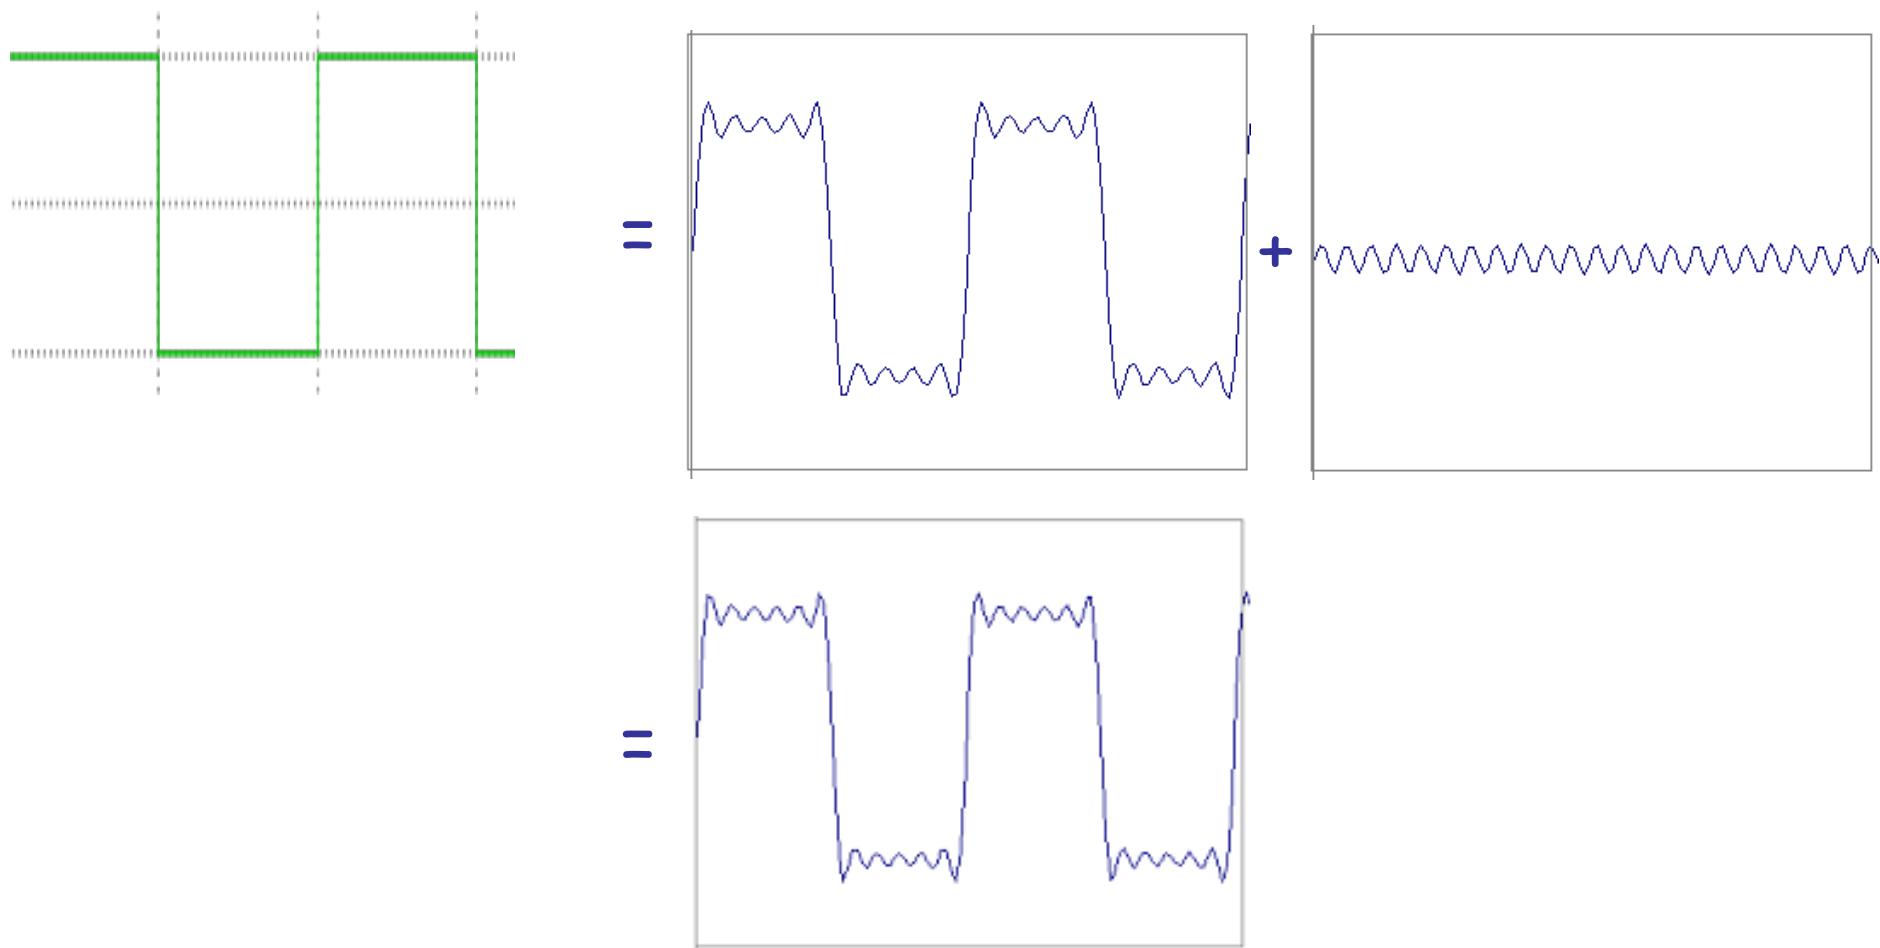
Slide credit: A. Efros

Frequency Spectra



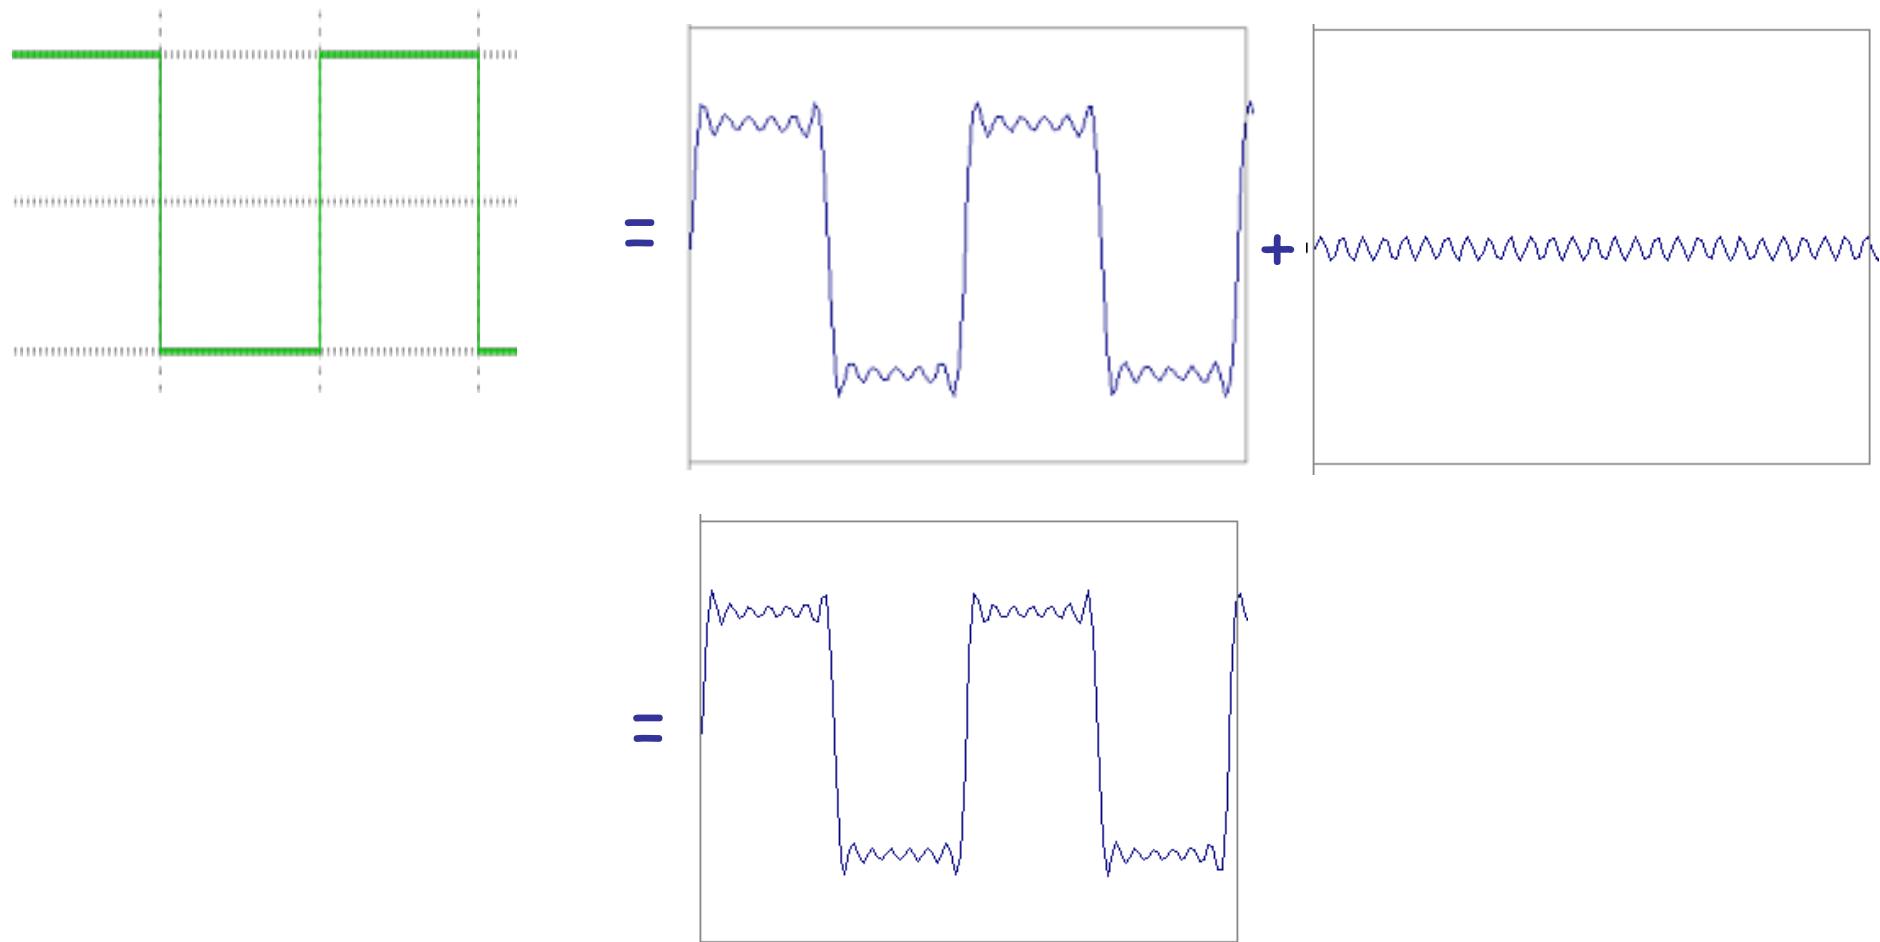
Slide credit: A. Efros

Frequency Spectra



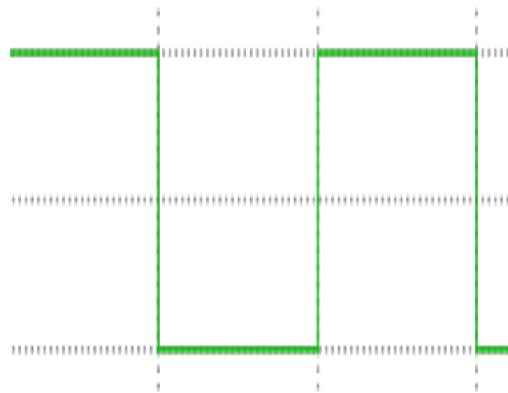
Slide credit: A. Efros

Frequency Spectra

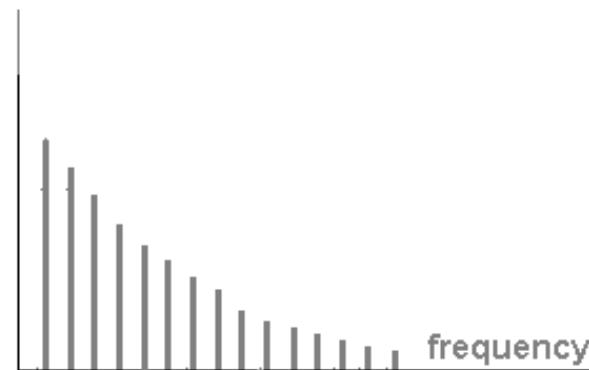


Slide credit: A. Efros

Frequency Spectra



$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



Frequency Spectra

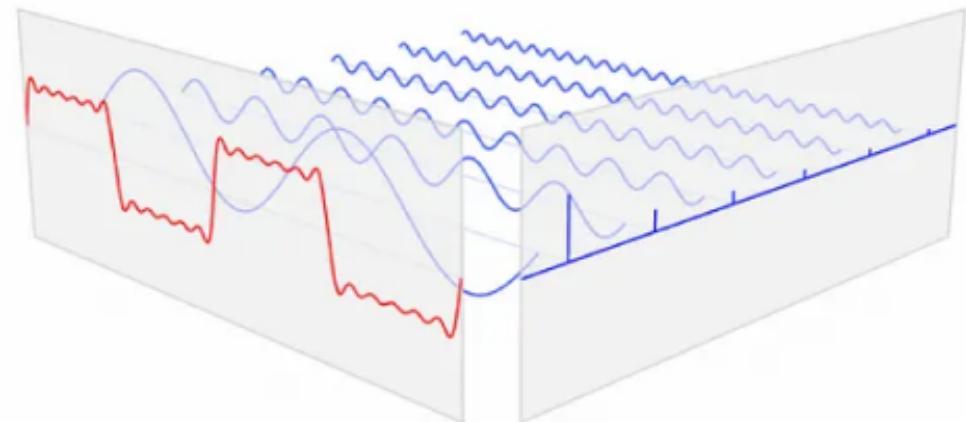
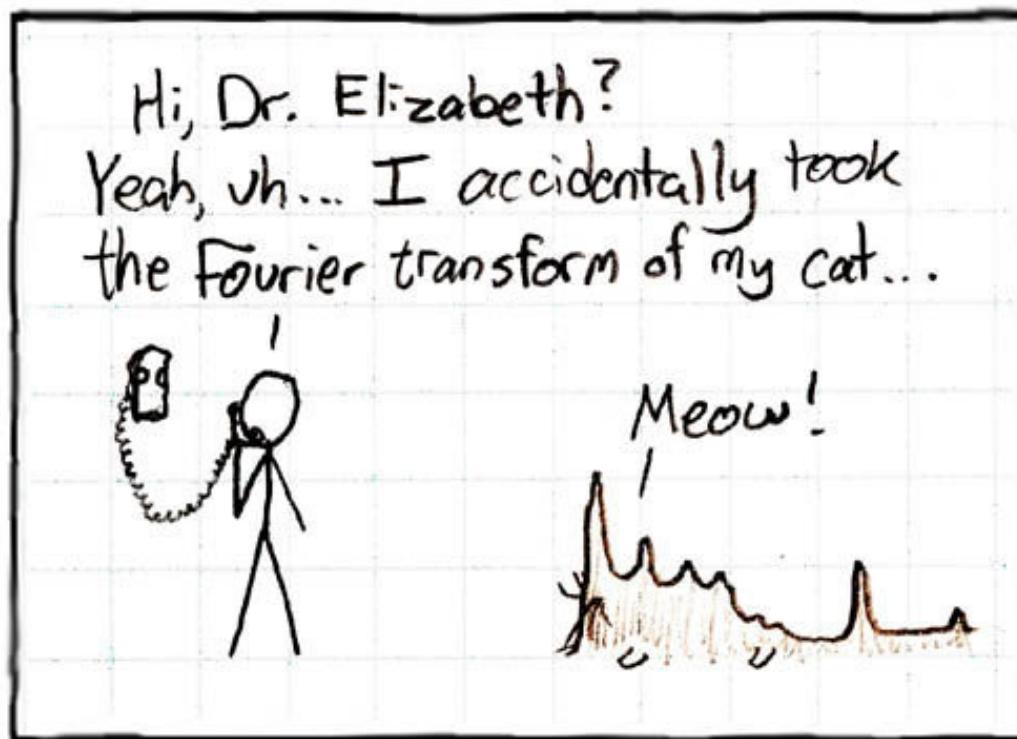


Image credit: Lucas V. Barbosa

Other signals

- We can also think of all kinds of other signals the same way



xkcd.com

Slide credit: J. Hays

Fourier Transform

We want to understand the frequency w of our signal. So, let's reparametrize the signal by w instead of x :



For every w from 0 to inf, $F(w)$ holds the amplitude A and phase f of the corresponding sine

$$A \sin(\omega x + \phi)$$

- How can F hold both? Complex number trick!

$$\begin{aligned} F(\omega) &= R(\omega) + iI(\omega) \\ A &= \pm \sqrt{R(\omega)^2 + I(\omega)^2} \quad \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)} \end{aligned}$$

We can always go back:



Fourier Transform

- Fourier transform stores the magnitude and phase at each frequency
 - Magnitude encodes how much signal there is at a particular frequency
 - Phase encodes spatial information (indirectly)
 - For mathematical convenience, this is often notated in terms of real and complex numbers

$$\text{Amplitude: } A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$\text{Phase: } \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

Discrete Fourier transform

- Forward transform

$$F(u,v) = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x,y) e^{-j2\pi(ux/M+vy/N)}$$

for $u = 0, 1, 2, \dots, M - 1, v = 0, 1, 2, \dots, N - 1$

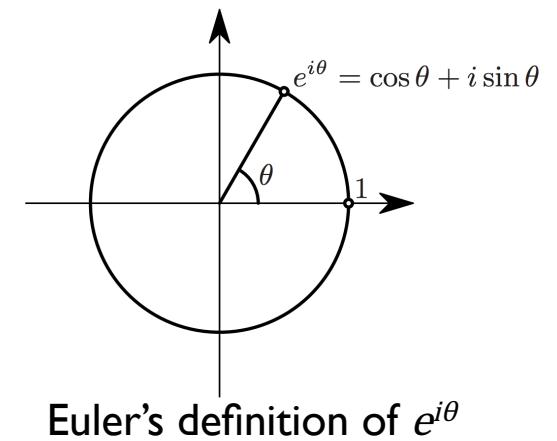
- Inverse transform

$$f(x,y) = \sum_{u=0}^{M-1} \sum_{v=0}^{N-1} F(u,v) e^{j2\pi(ux/M+vy/N)}$$

for $x = 0, 1, 2, \dots, M - 1, y = 0, 1, 2, \dots, N - 1$

u, v : the transform or frequency variables

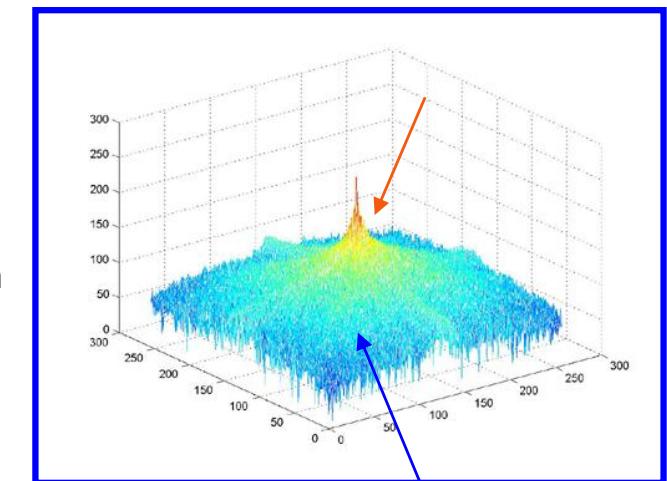
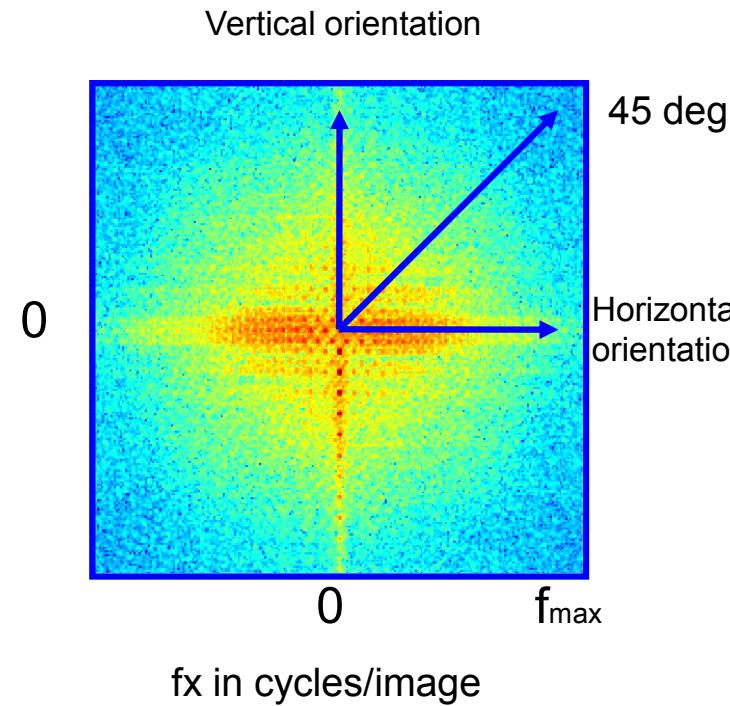
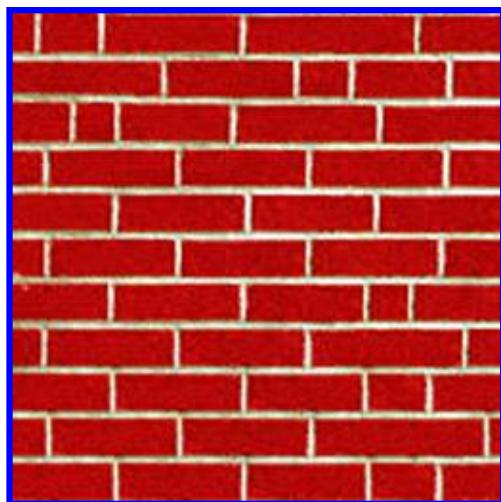
x, y : the spatial or image variables



The Fourier Transform

- Represent function on a new basis
 - Think of functions as vectors, with many components
 - We now apply a linear transformation to transform the basis
 - dot product with each basis element
- In the expression, u and v select the basis element, so a function of x and y becomes a function of u and v
- basis elements have the form $e^{-i2\pi(ux+vy)}$

How to interpret 2D Fourier Spectrum



**High
spatial
frequencies**

Log power spectrum

Fourier basis element

$$e^{-i2\pi(ux+vy)}$$

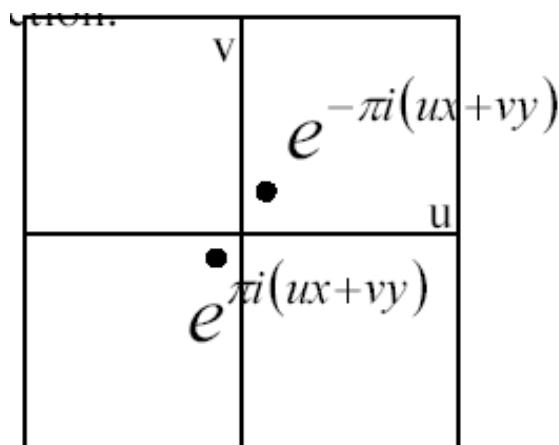
example, real part

$$F^{u,v}(x,y)$$

$F^{u,v}(x,y) = \text{const.}$ for $(ux+vy) = \text{const.}$

Vector (u,v)

- Magnitude gives frequency
- Direction gives orientation.



Slide credit: S. Thrun

Here u and v are larger than in the previous slide.

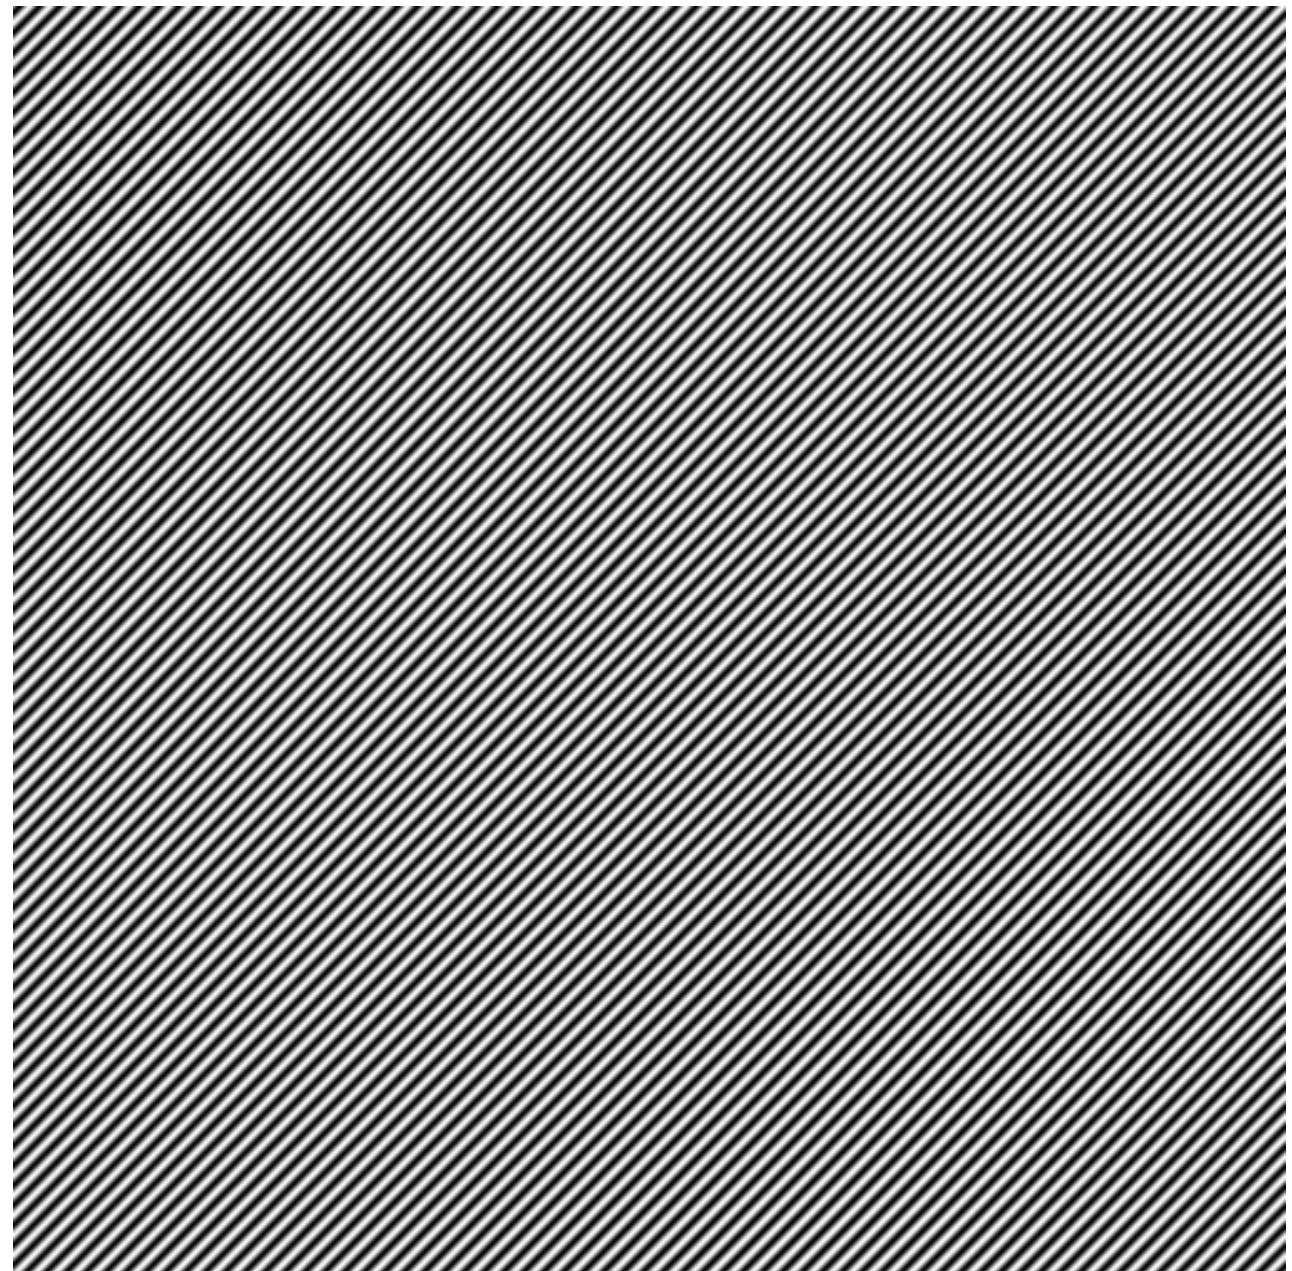
$$\begin{array}{|c|c|} \hline e^{-\pi i \frac{v}{u}(ux+vy)} & \\ \hline \bullet & u \\ \hline & \bullet e^{\pi i \frac{v}{u}(ux+vy)} \\ \hline \end{array}$$



And larger still...

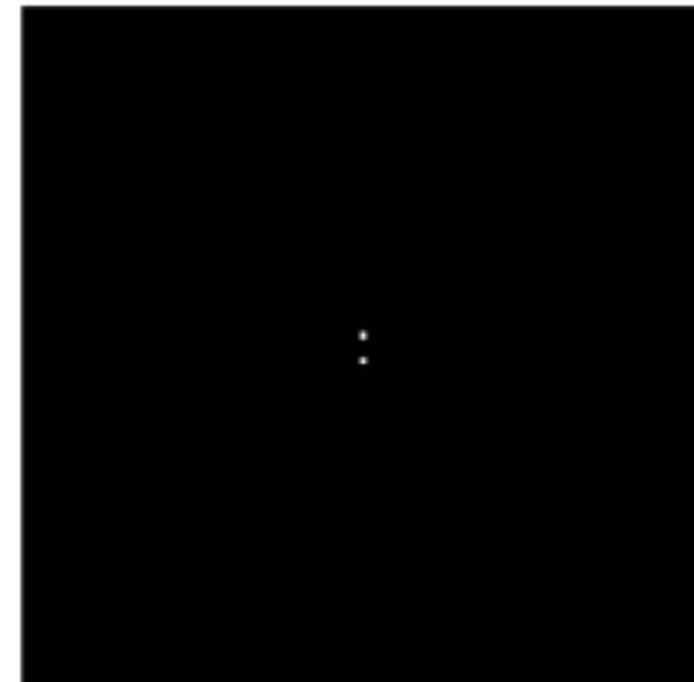
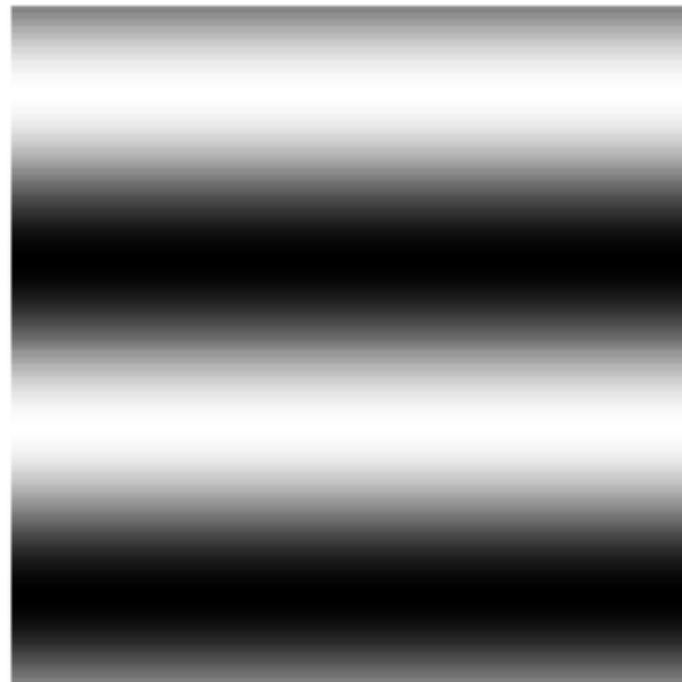
$$e^{-\pi i(ux+vy)}$$
$$e^{\pi i(ux+vy)}$$

A 2x2 grid divided into four quadrants. The top-left quadrant contains a black dot labeled 'v'. The bottom-right quadrant contains a black dot labeled 'u'.



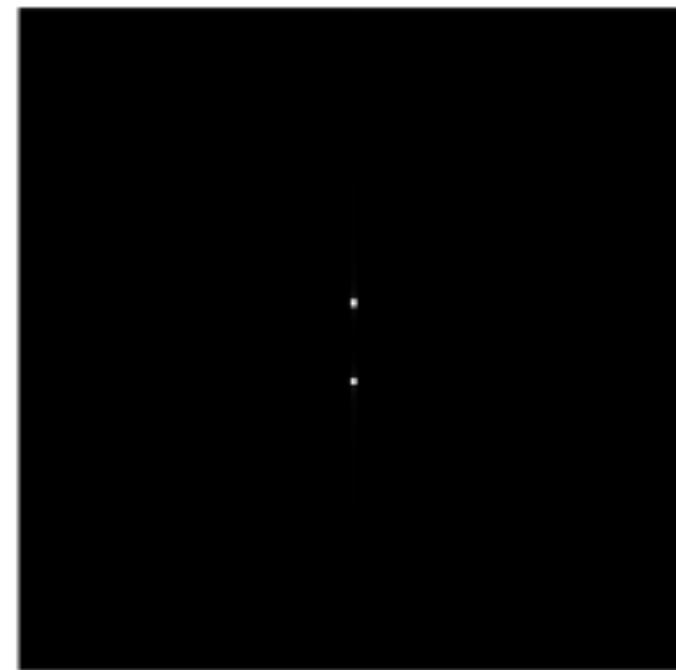
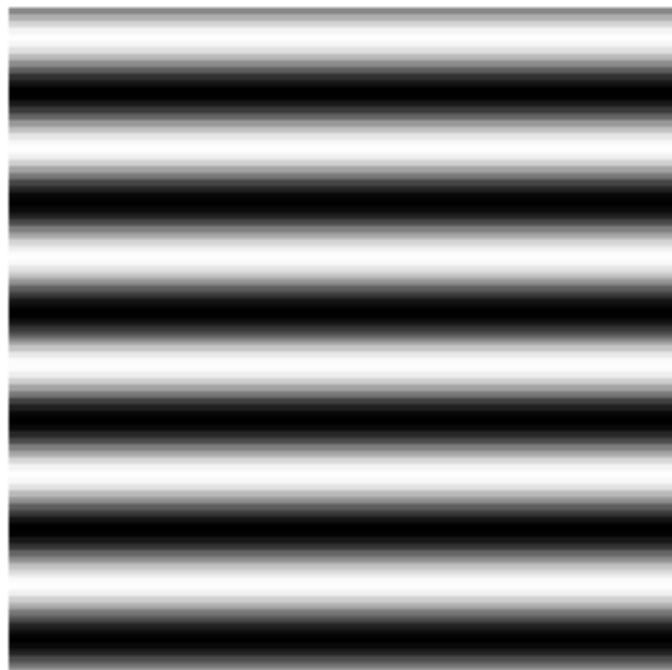
Slide credit: S. Thrun

2D FFT



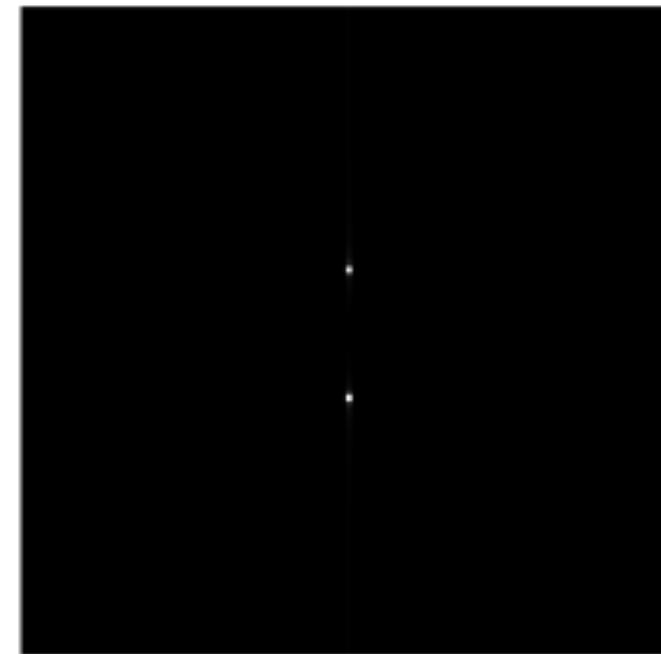
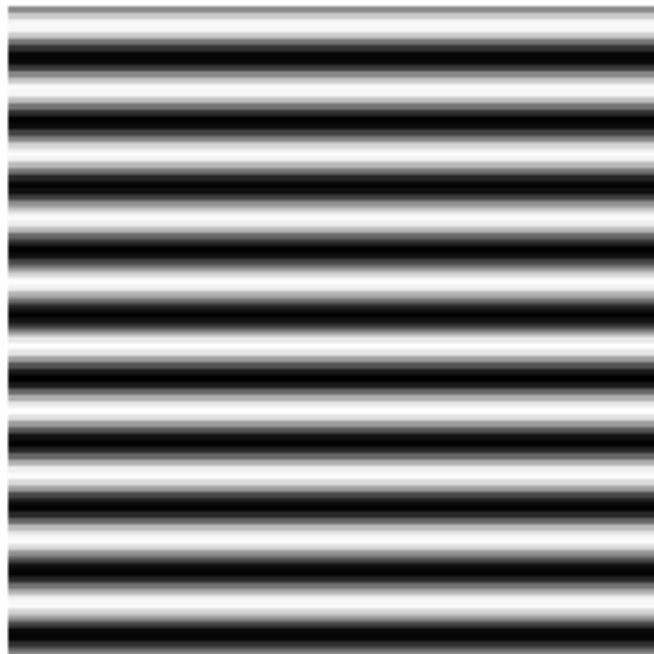
Sinusoid with frequency = 1 and its FFT

2D FFT



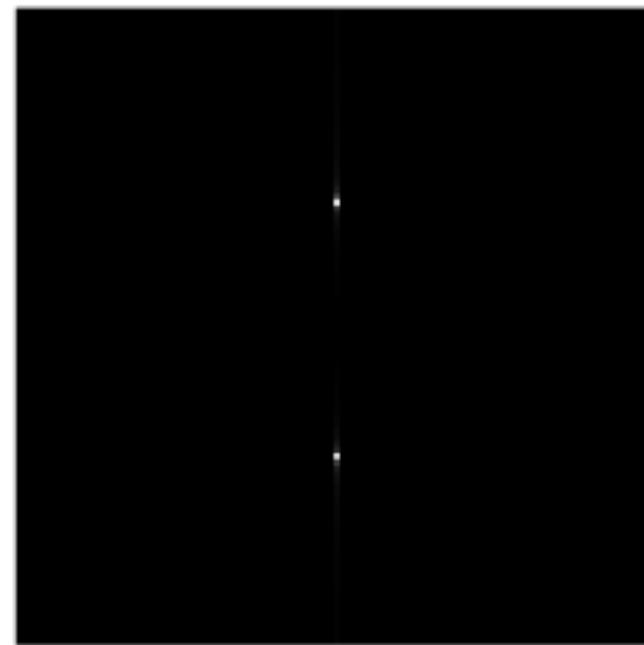
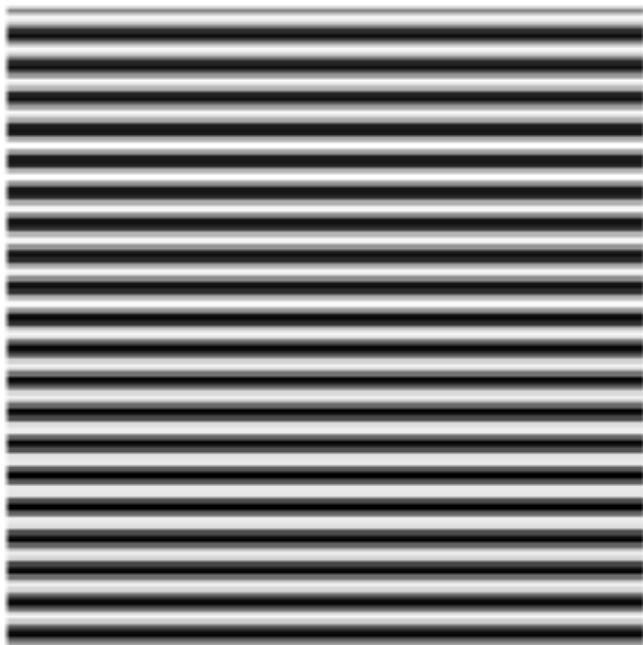
Sinusoid with frequency = 3 and its FFT

2D FFT



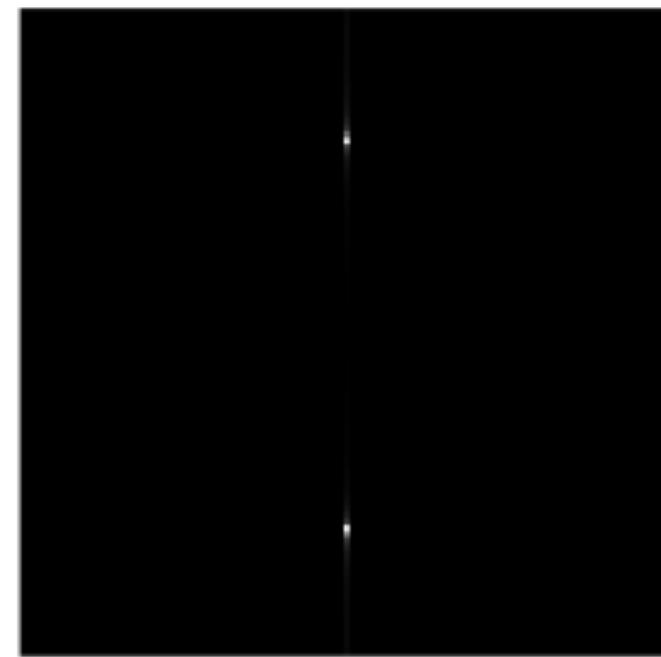
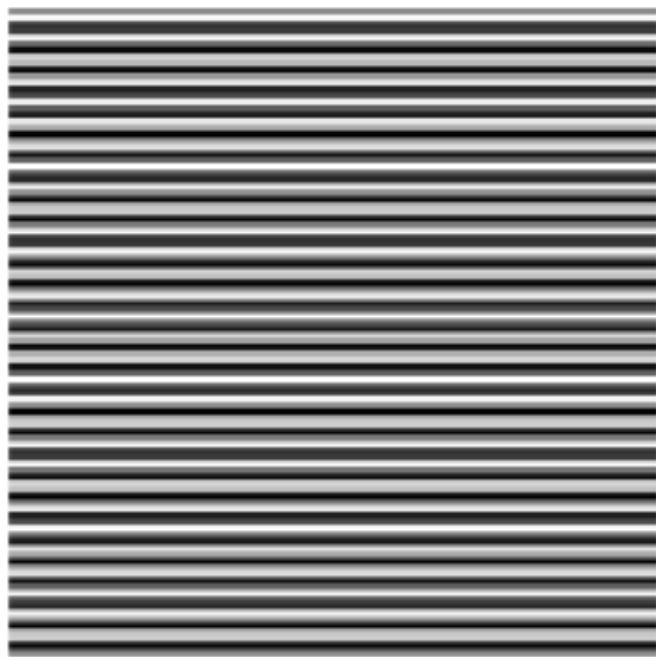
Sinusoid with frequency = 5 and its FFT

2D FFT



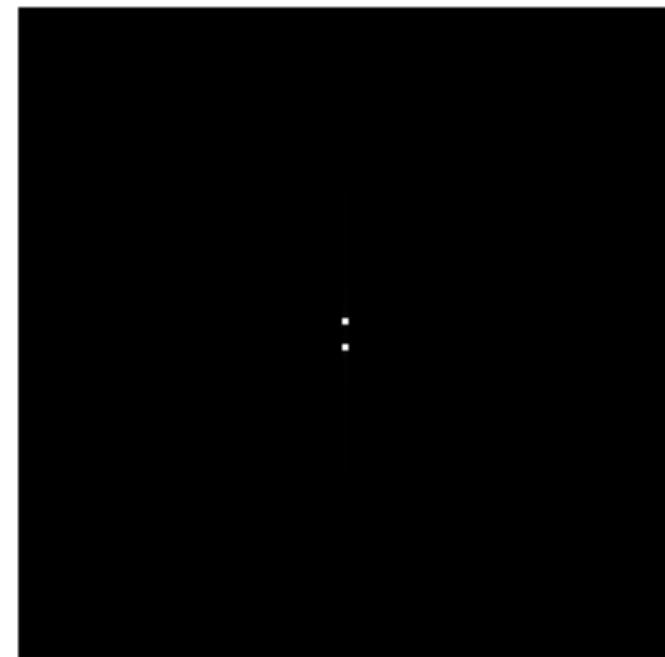
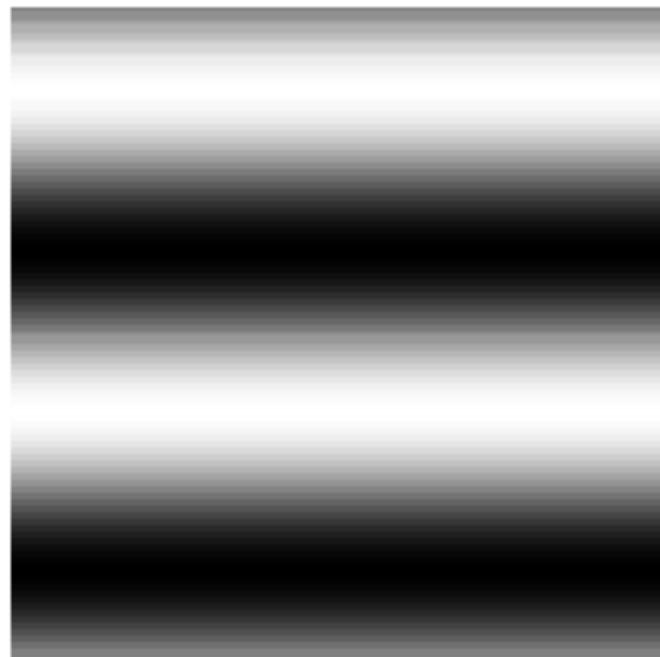
Sinusoid with frequency = 10 and its FFT

2D FFT



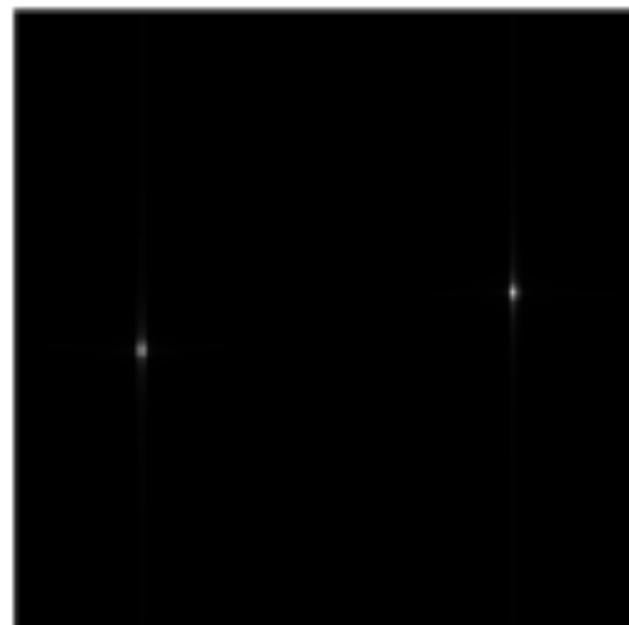
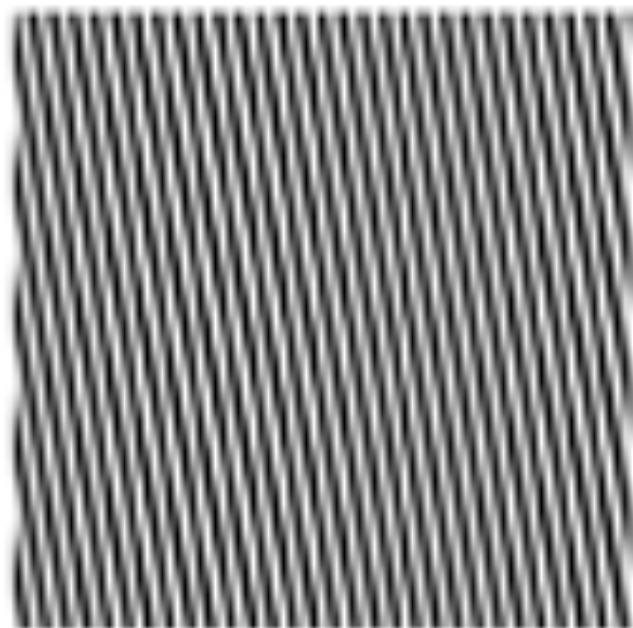
Sinusoid with frequency = 15 and its FFT

2D FFT



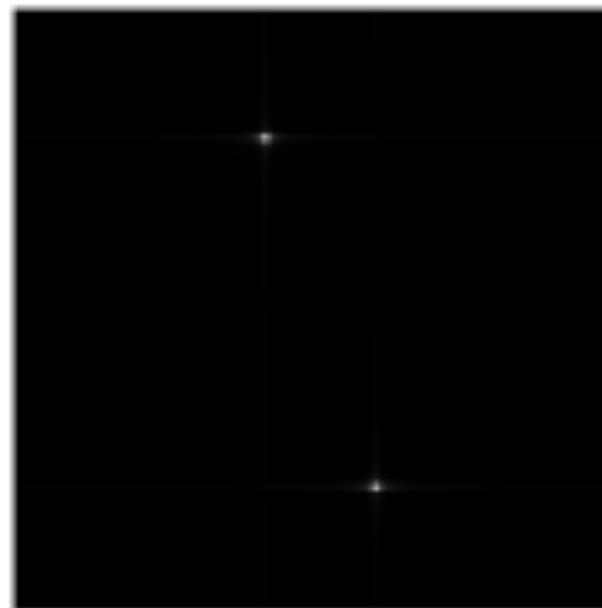
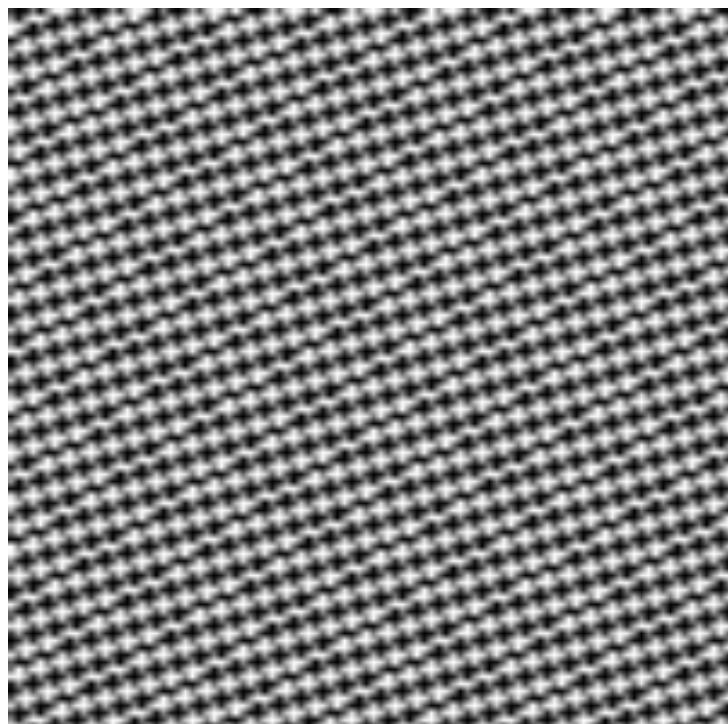
Sinusoid with varying frequency and their FFT

Rotation



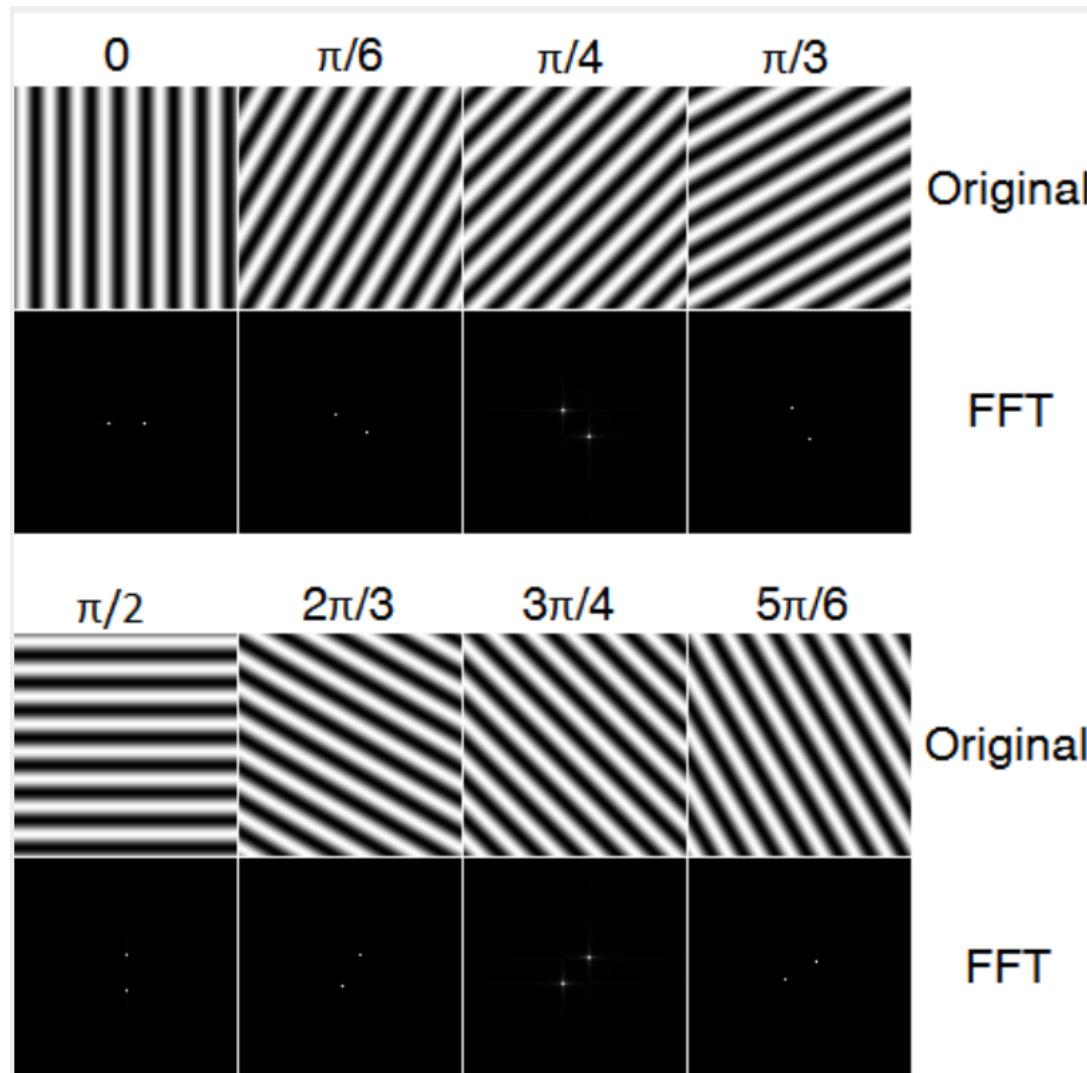
Sinusoid rotated at 30 degrees and its FFT

2D FFT



Sinusoid rotated at 60 degrees and its FFT

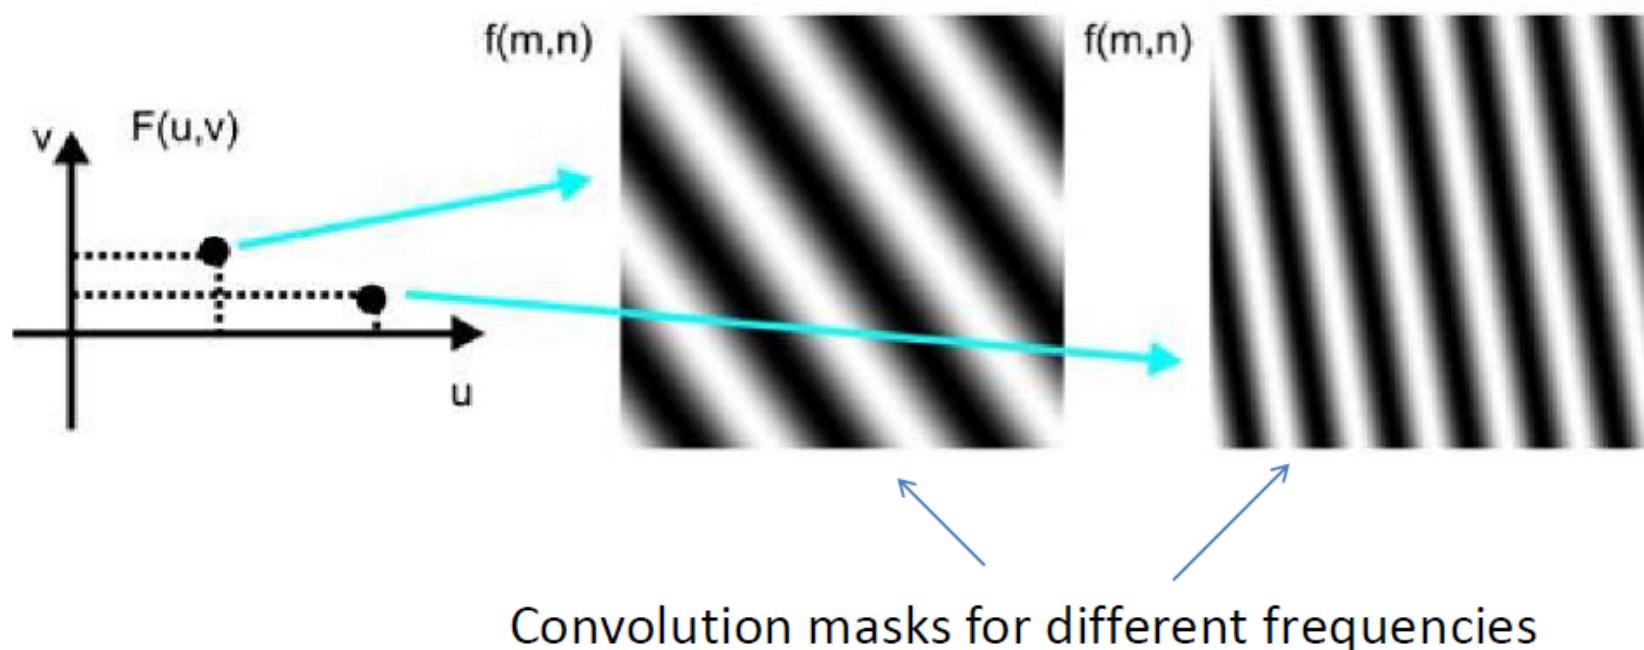
2D FFT



Rotation of sinusoid
from 0 to π
and their FFTs.

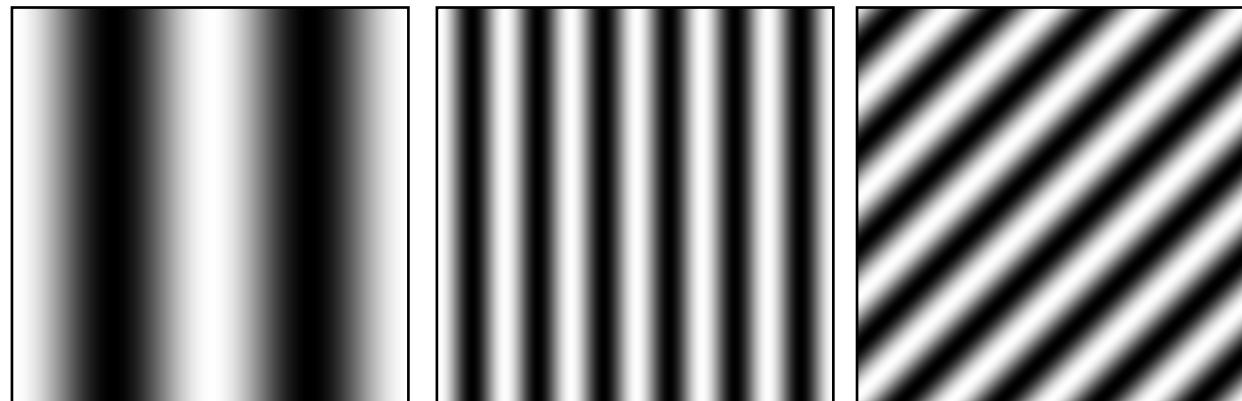
2D FFT

$$F(u, v) = \frac{1}{MN} \cdot \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-i2\pi(xu/M + yv/N)}$$



Fourier analysis in images

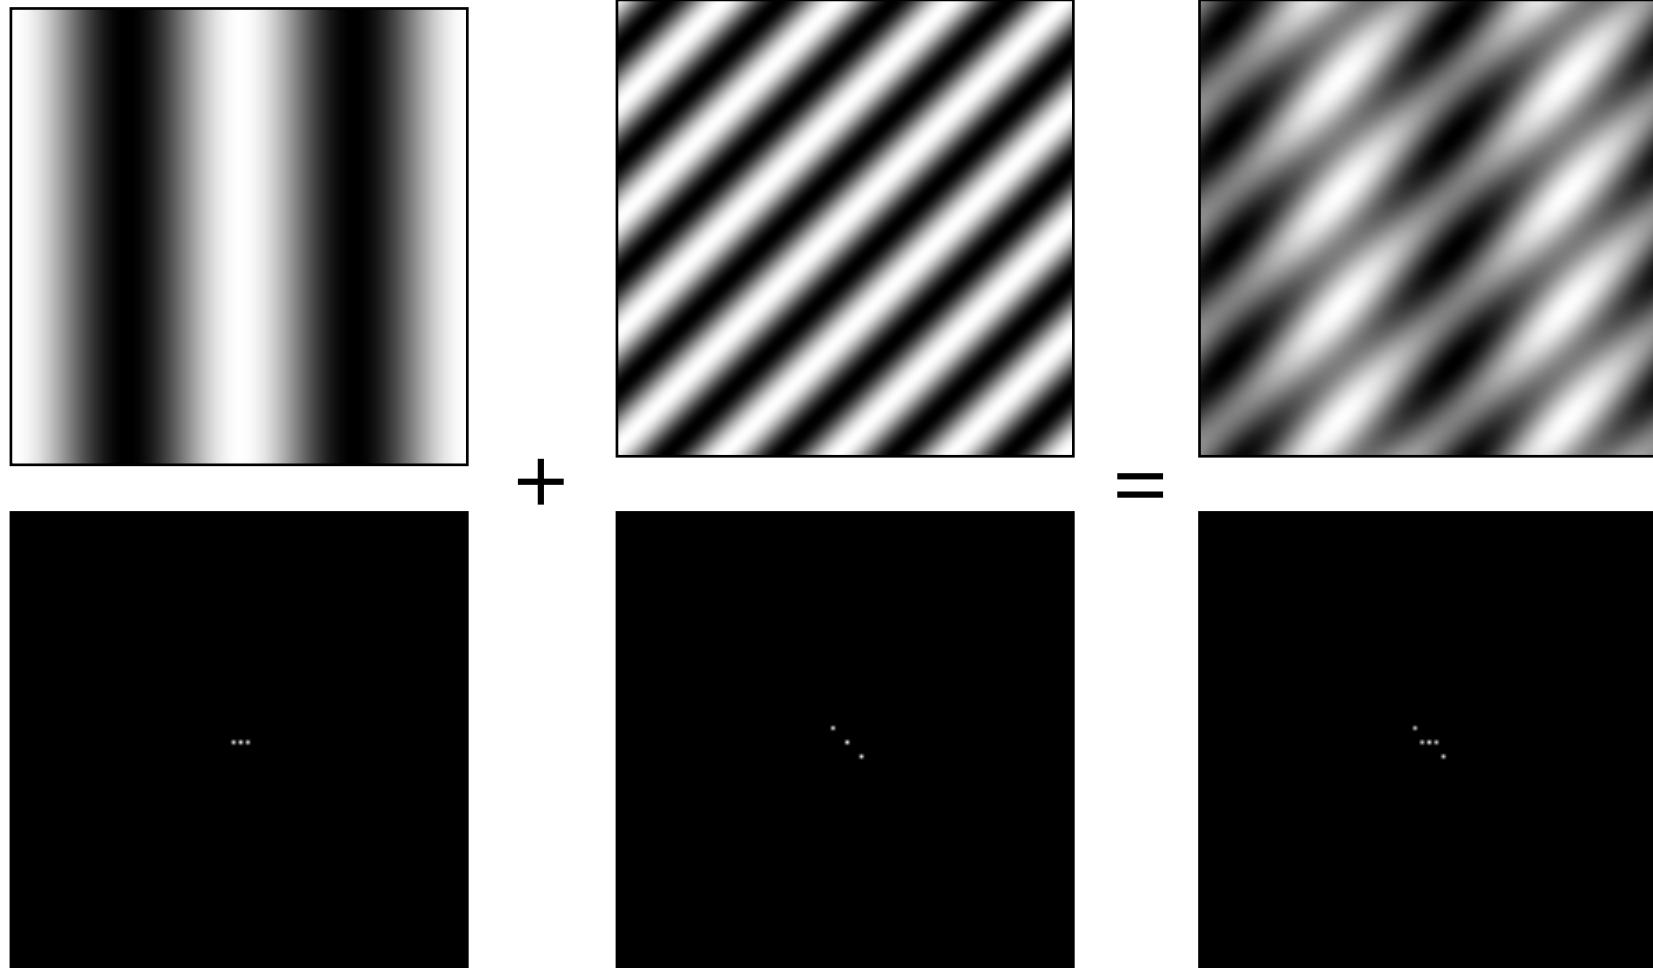
Intensity Image



Fourier Image



Signals can be composed

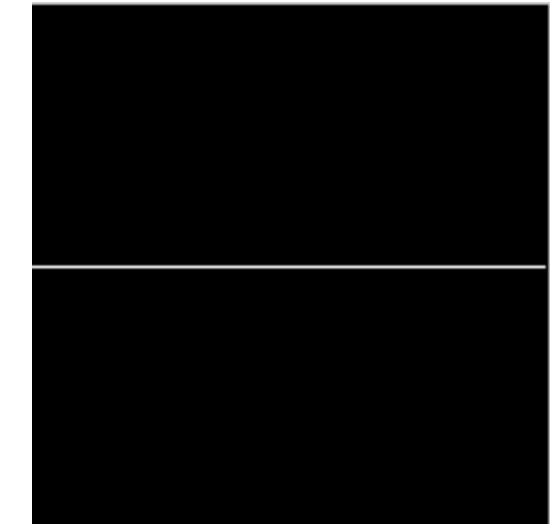
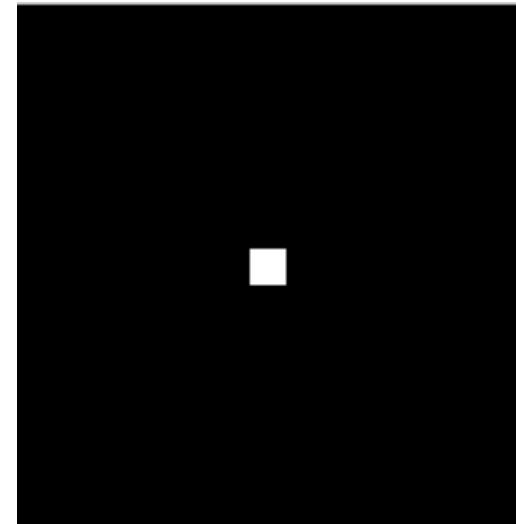
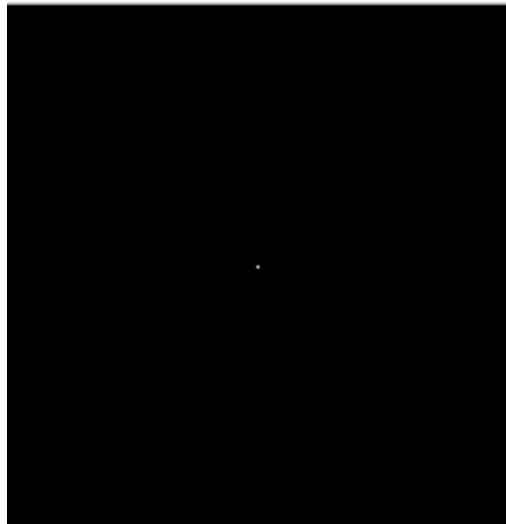


Slide credit: A. Efros

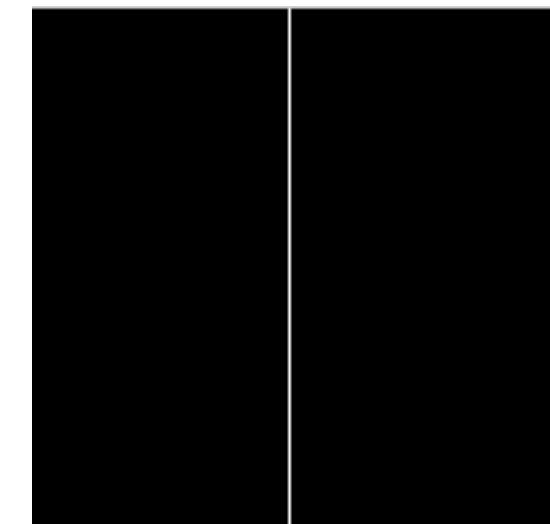
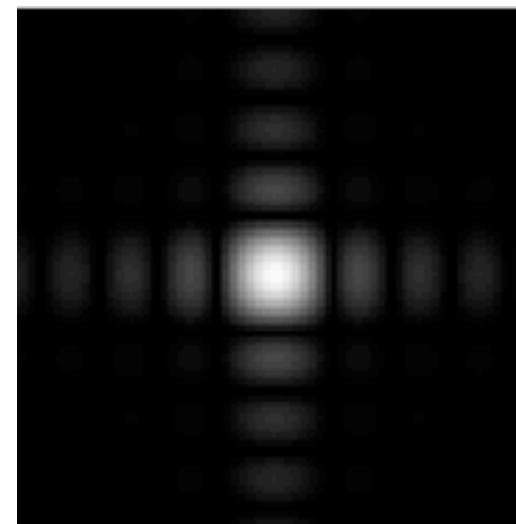
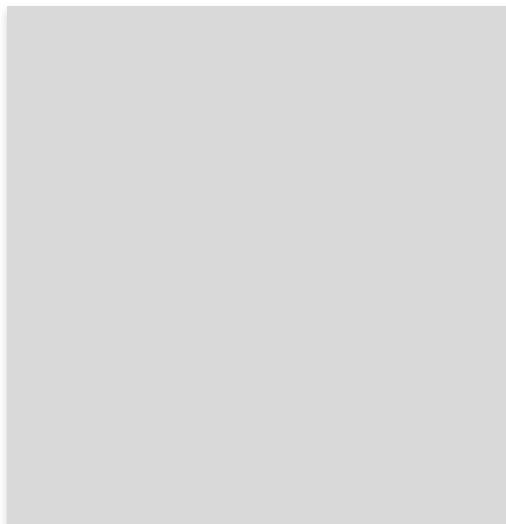
<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

Some important Fourier Transforms

Image



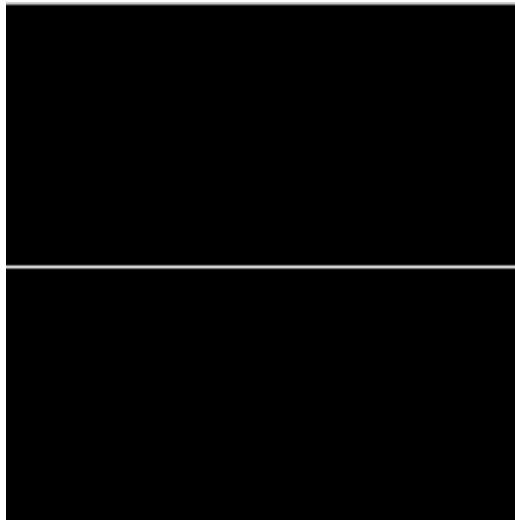
Magnitude FT



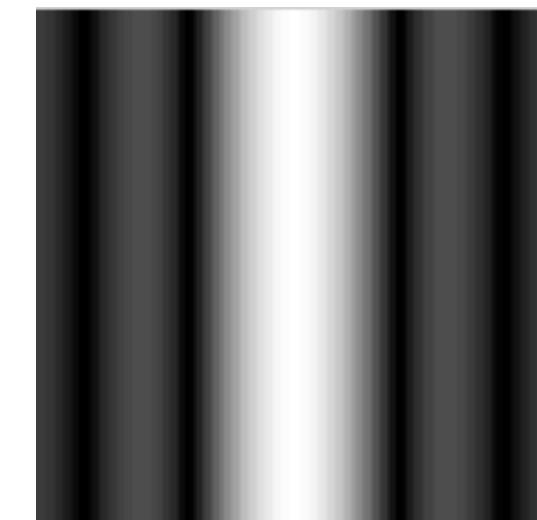
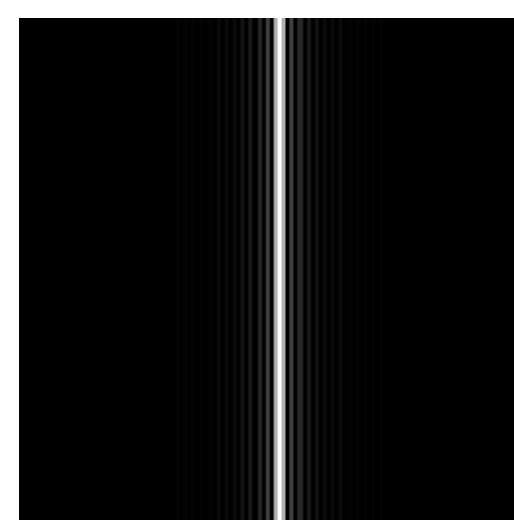
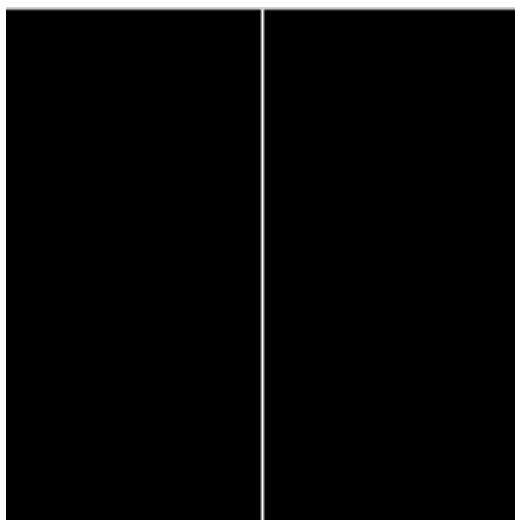
Slide credit: B. Freeman and A. Torralba

Some important Fourier Transforms

Image



Magnitude FT



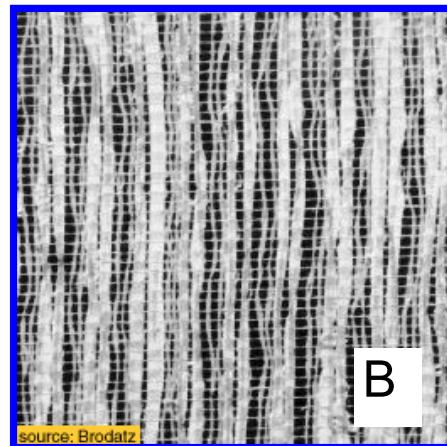
Slide credit: B. Freeman and A. Torralba

Fourier Amplitude Spectrum

Images



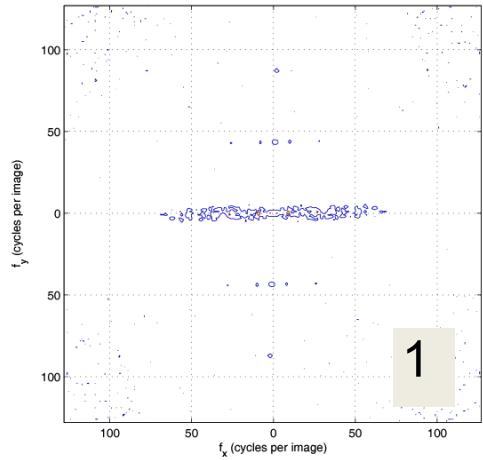
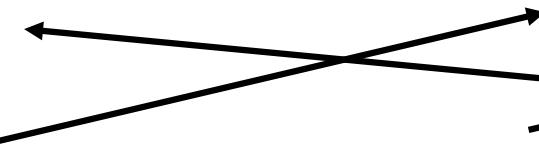
A



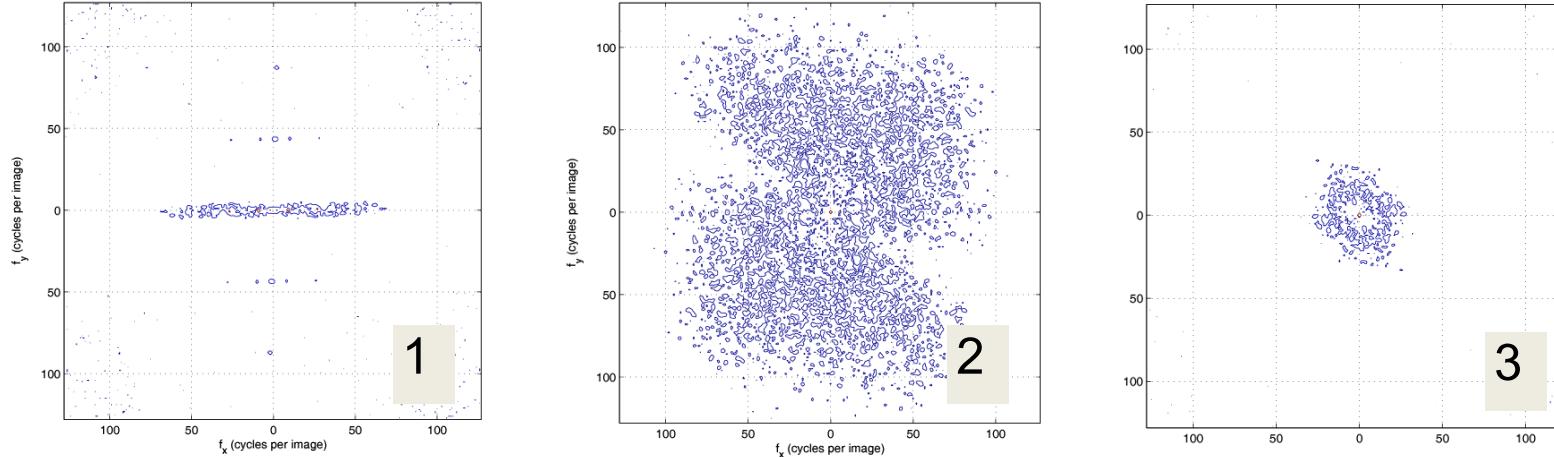
B



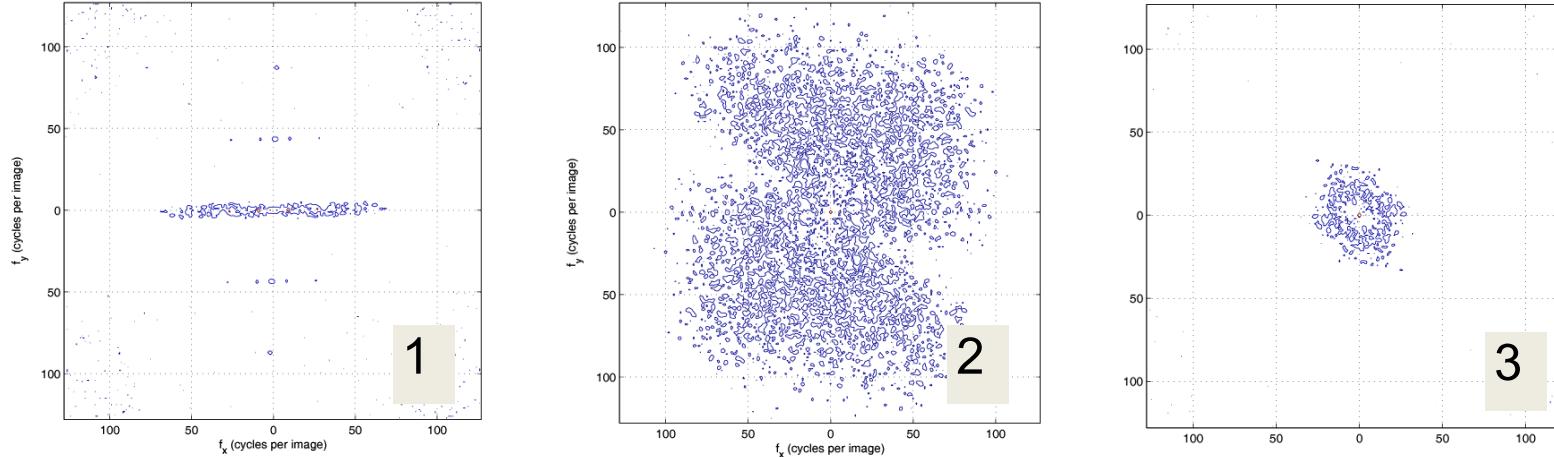
C



1



2



3

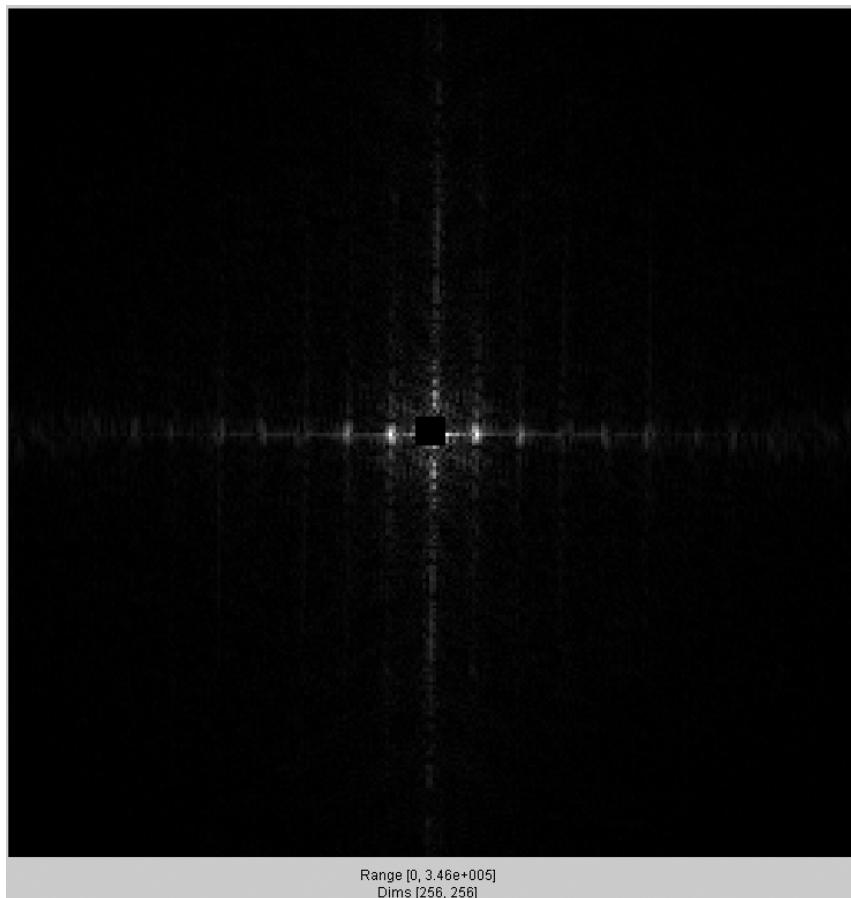
f_x (cycles/image pixel size)

f_x (cycles/image pixel size)

f_x (cycles/image pixel size)

Slide credit: B. Freeman and A. Torralba

Fourier transform magnitude



Slide credit: B. Freeman and A. Torralba

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

$$\mathcal{F}^{-1}[gh] = \mathcal{F}^{-1}[g] * \mathcal{F}^{-1}[h]$$

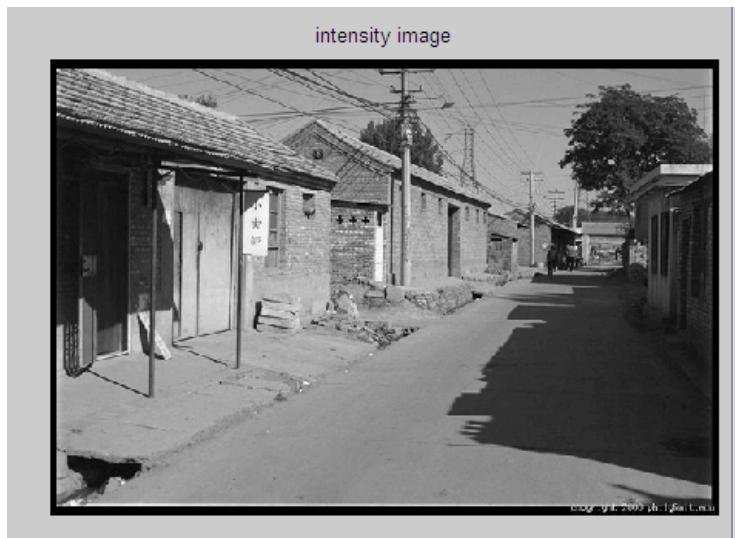
- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

Properties of Fourier Transforms

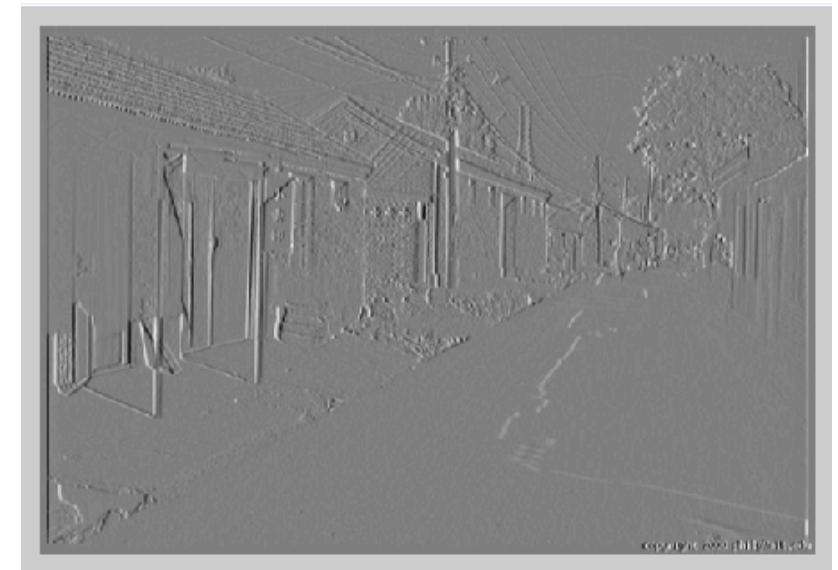
- **Linearity** $\mathcal{F}[ax(t) + by(t)] = a\mathcal{F}[x(t)] + b\mathcal{F}[y(t)]$
- Fourier transform of a real signal is symmetric about the origin
- The energy of the signal is the same as the energy of its Fourier transform

Filtering in spatial domain

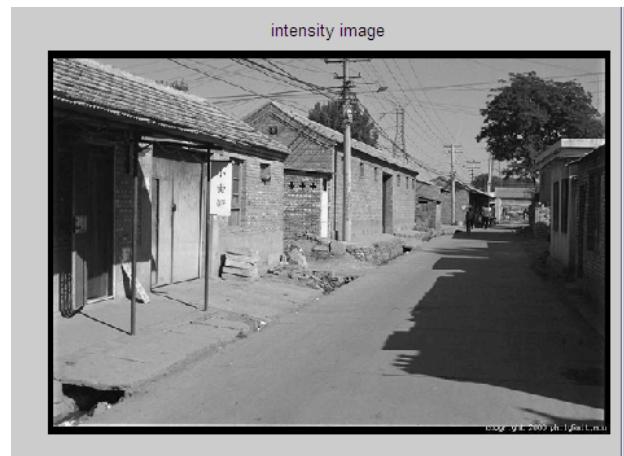
1	0	-1
2	0	-2
1	0	-1



$$\ast \quad \begin{matrix} \blacksquare & \square & \square \\ \square & \square & \square \\ \square & \square & \square \end{matrix} =$$

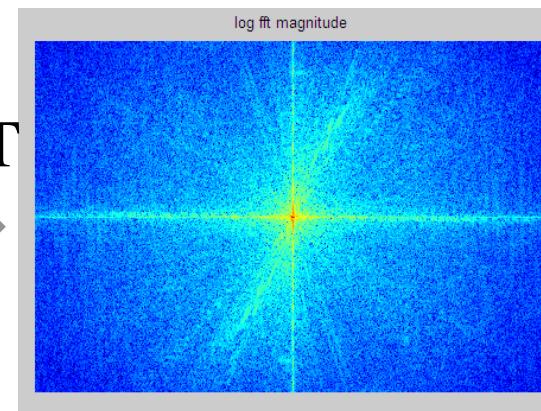


Filtering in frequency domain

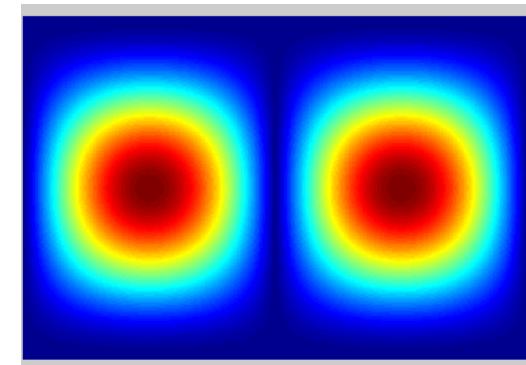


FFT

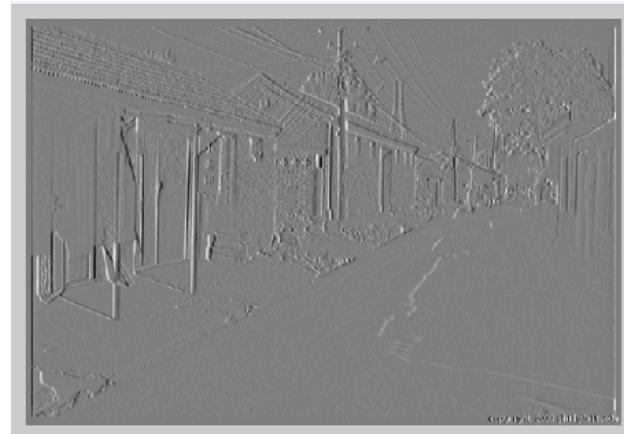
A gray arrow points from the intensity image to the next stage in the process, which is the log FFT magnitude image.



X

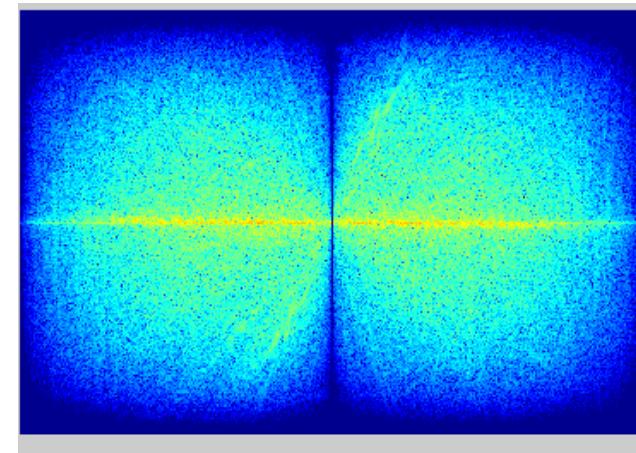


||

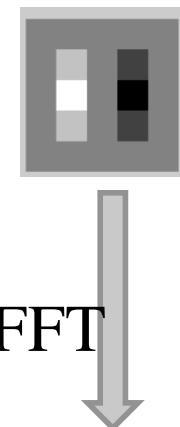


Inverse FFT

A gray arrow points from the filtered log FFT magnitude image back to the intensity image, indicating the inverse process of filtering.



Slide credit: D. Hoiem



FFT

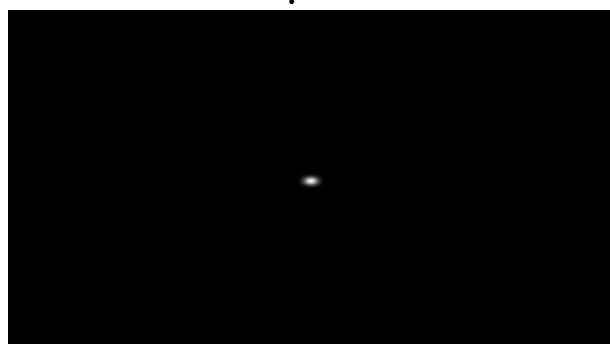
A gray arrow points from the filter diagram down to the color bar, indicating the flow of the process.

2D convolution theorem example

$f(x,y)$



$h(x,y)$



$g(x,y)$



*

\times

$|F(s_x, s_y)|$

$|H(s_x, s_y)|$

\Downarrow

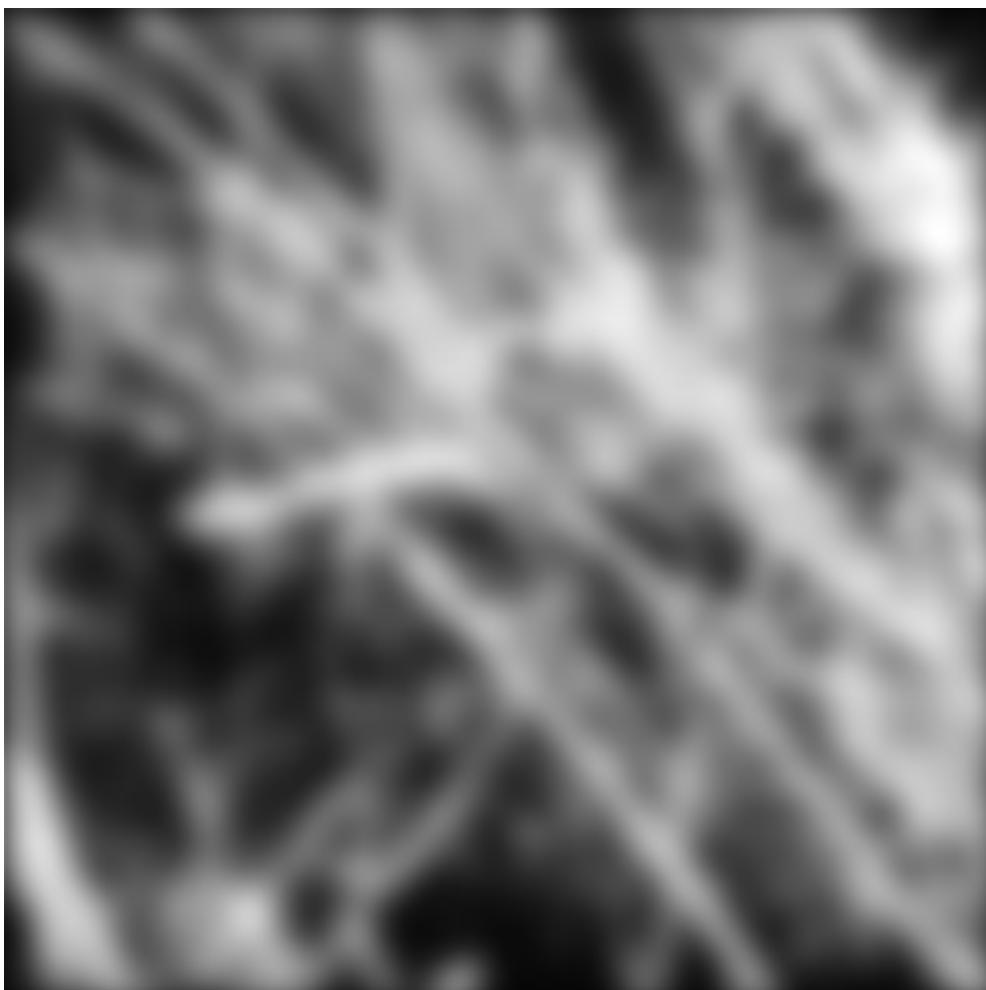
$|G(s_x, s_y)|$

Slide credit: A. Efros

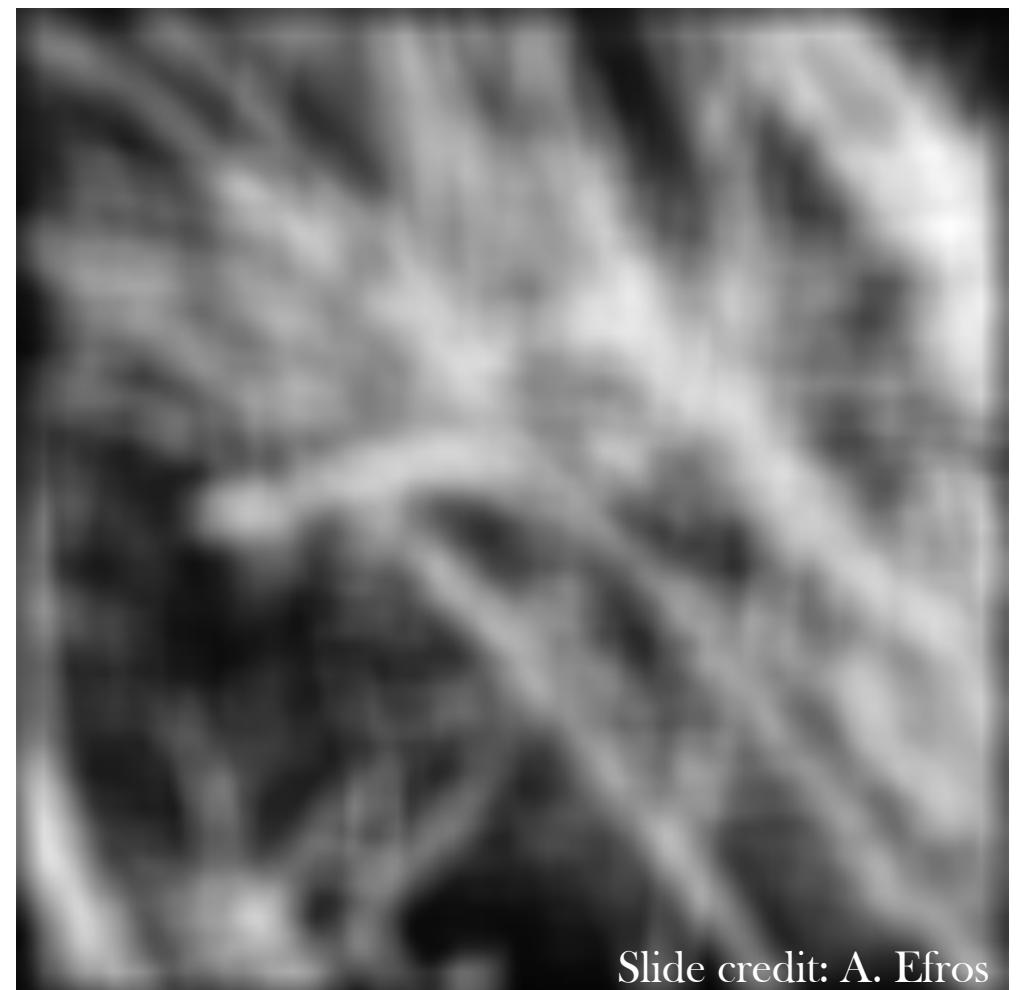
Filtering

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

Gaussian



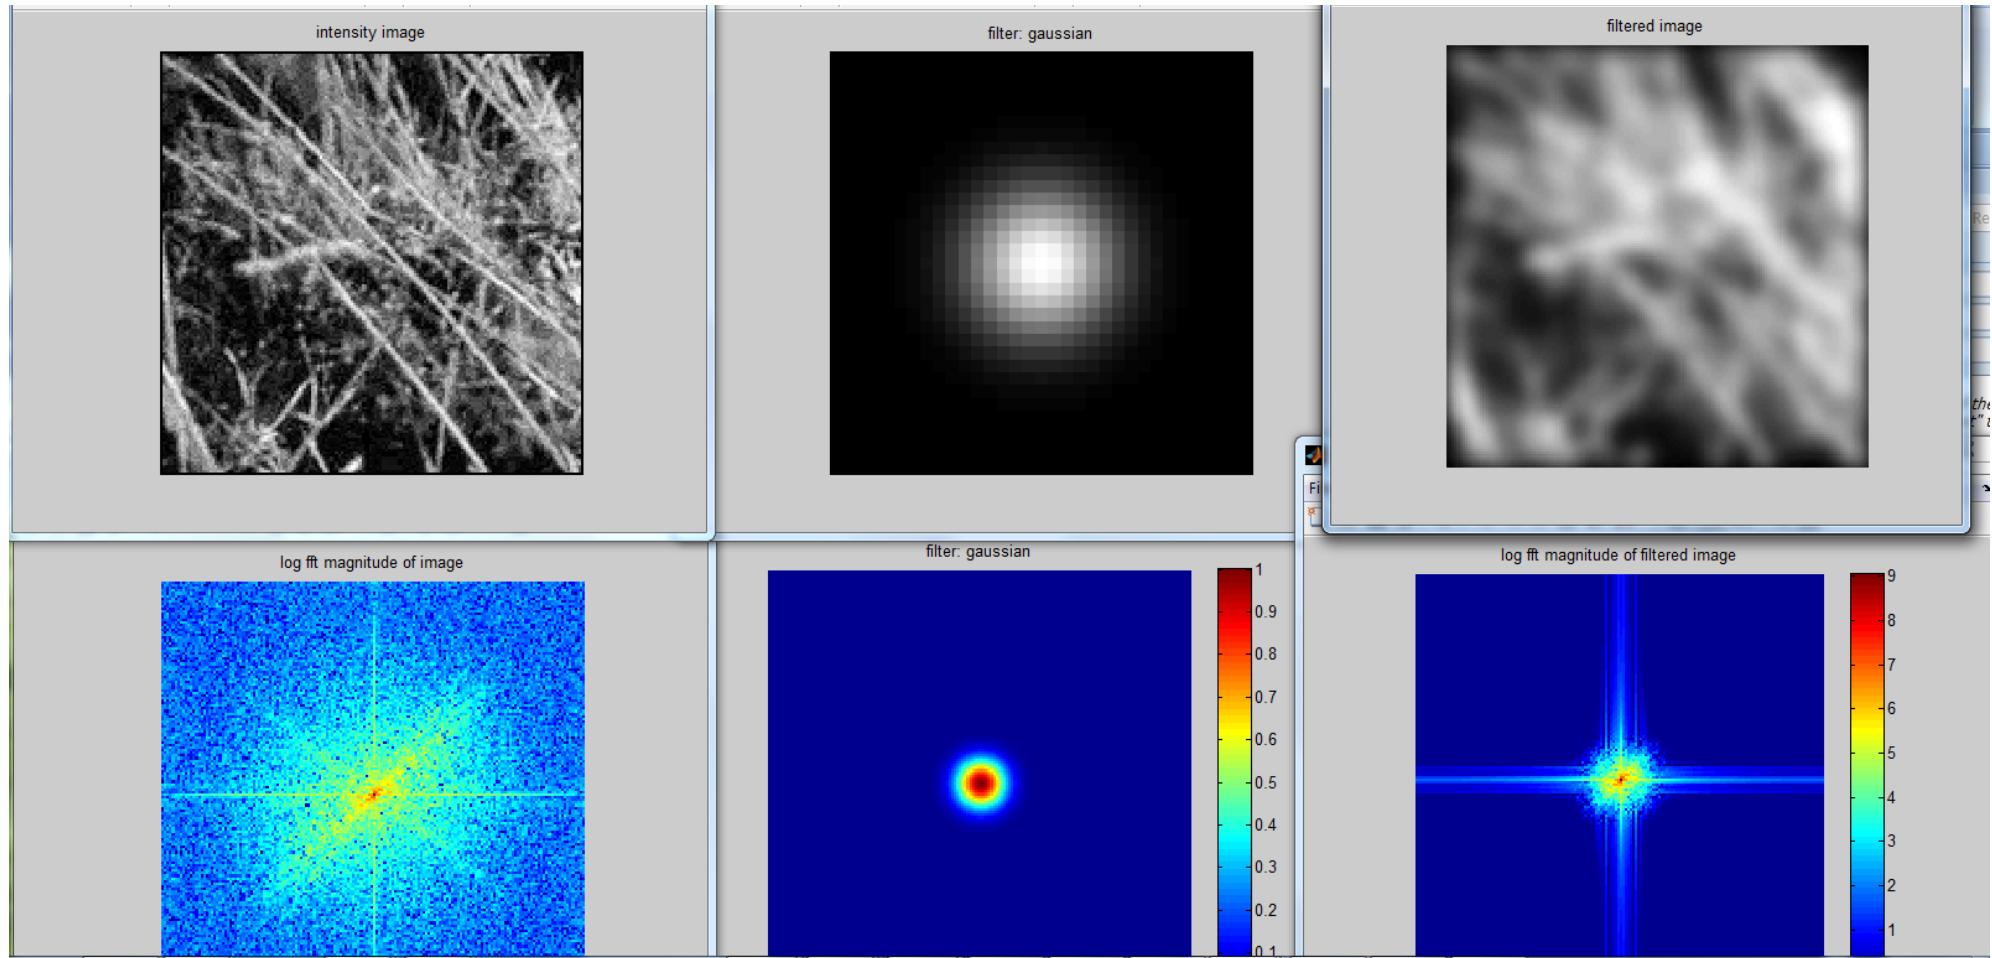
Box filter



Slide credit: A. Efros

Filtering

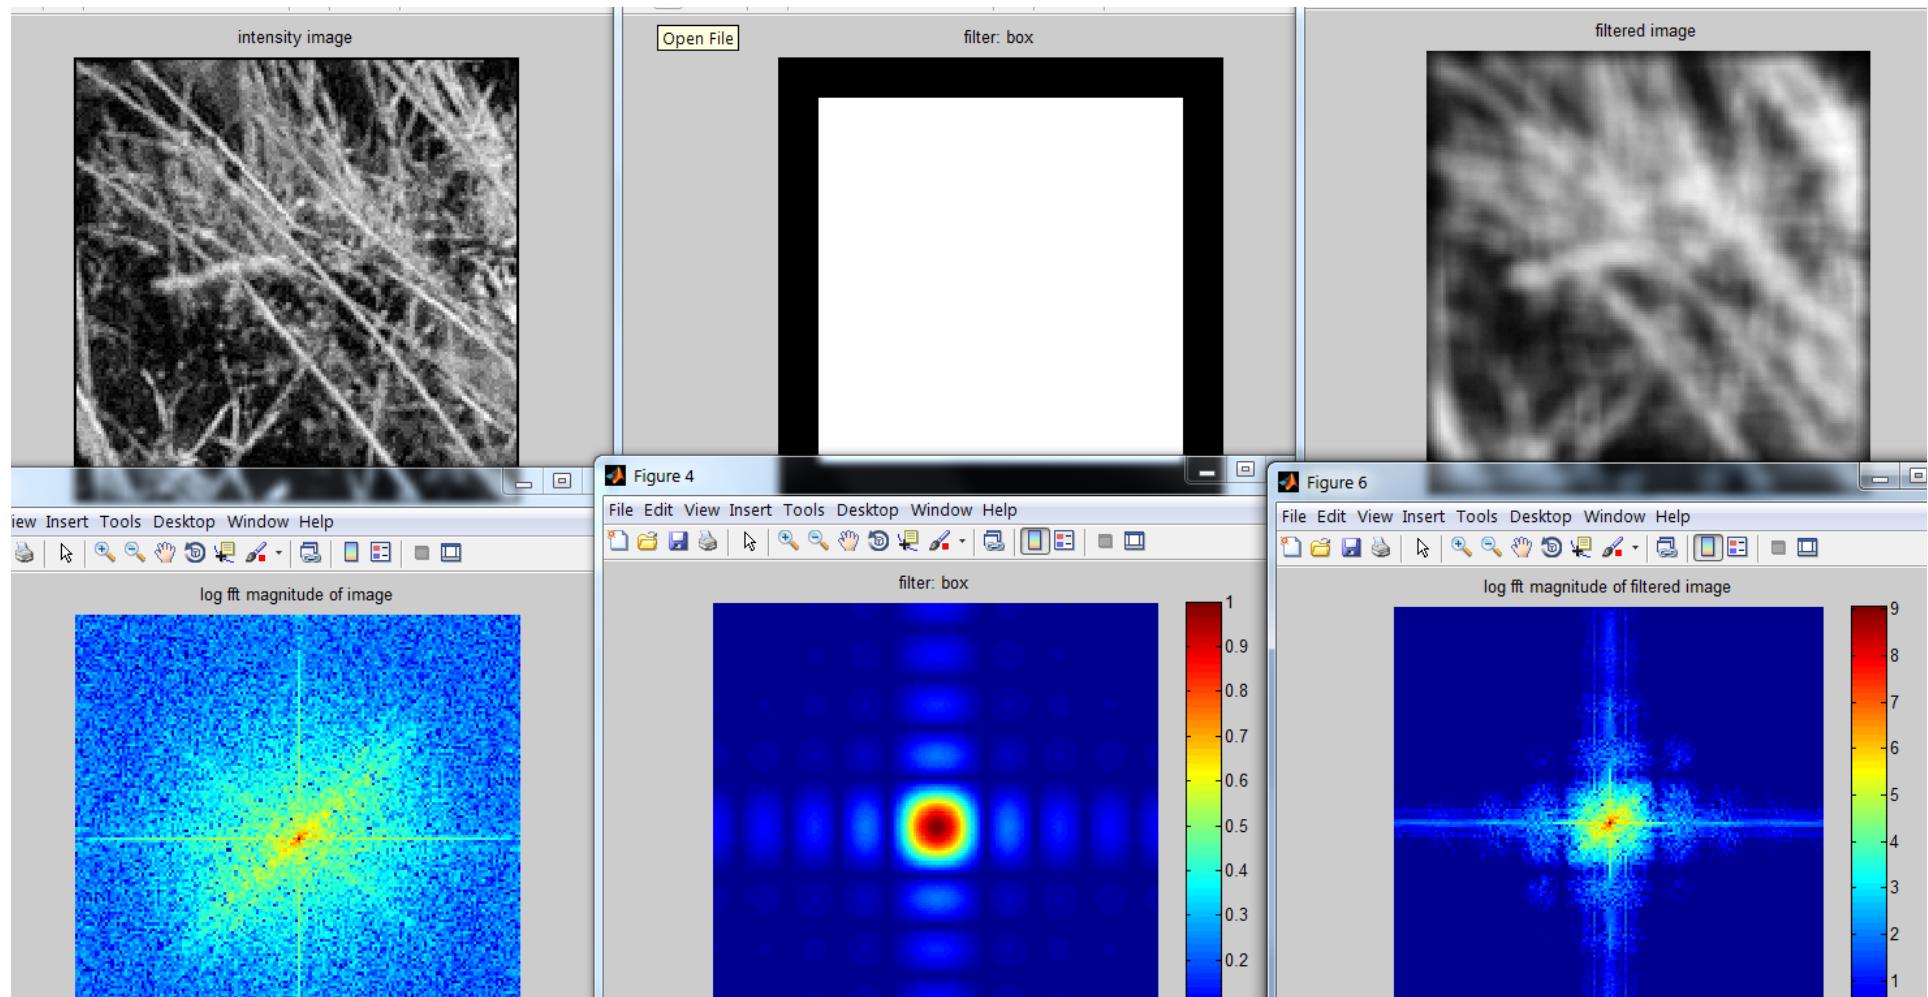
Gaussian



Slide credit: A. Efros

Filtering

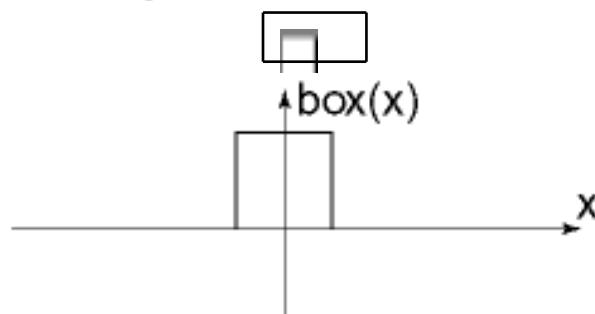
Box Filter



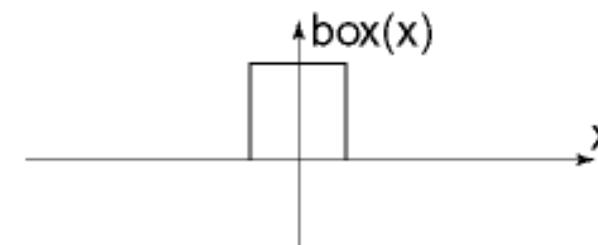
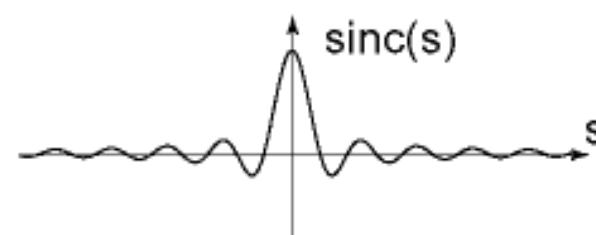
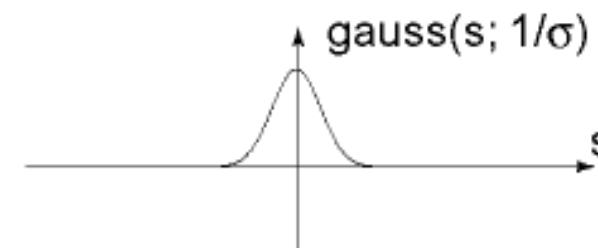
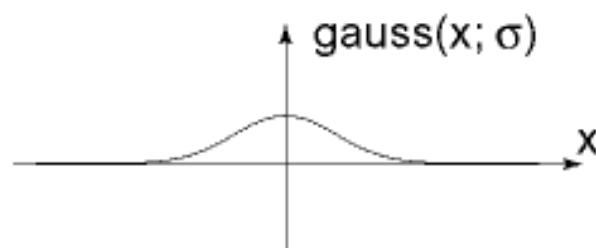
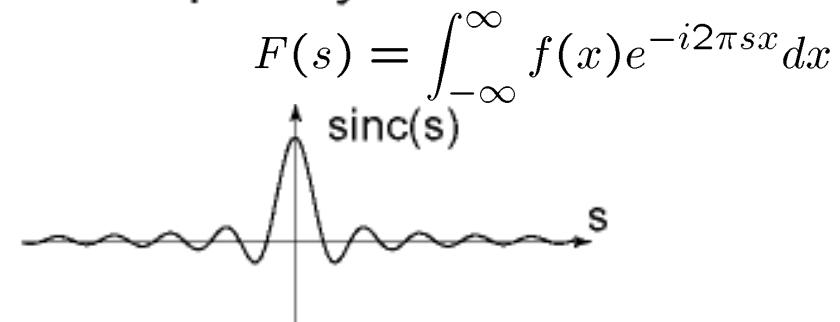
Slide credit: A. Efros

Fourier Transform pairs

Spatial domain

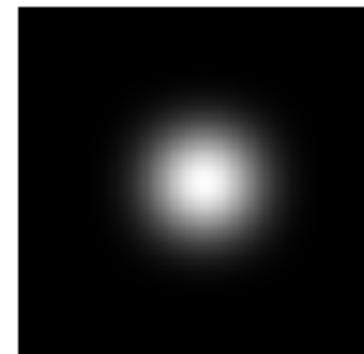
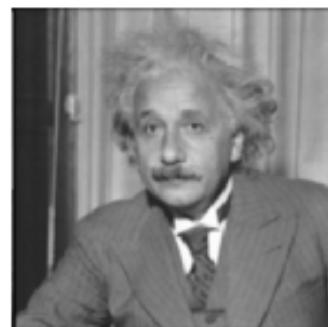
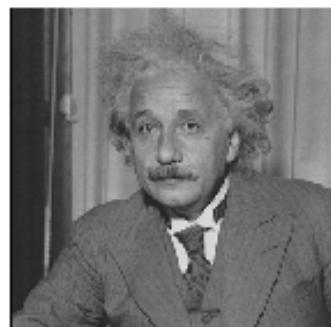


Frequency domain

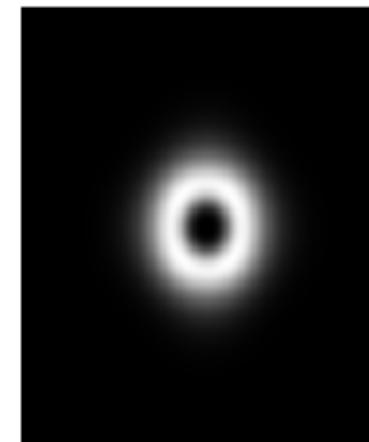
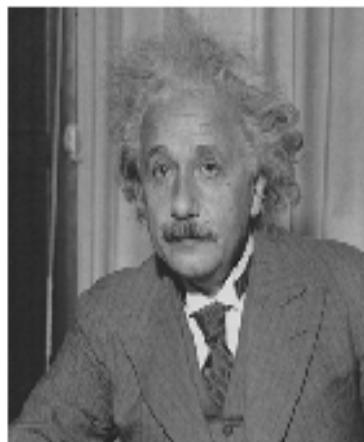


Low-pass, Band-pass, High-pass filters

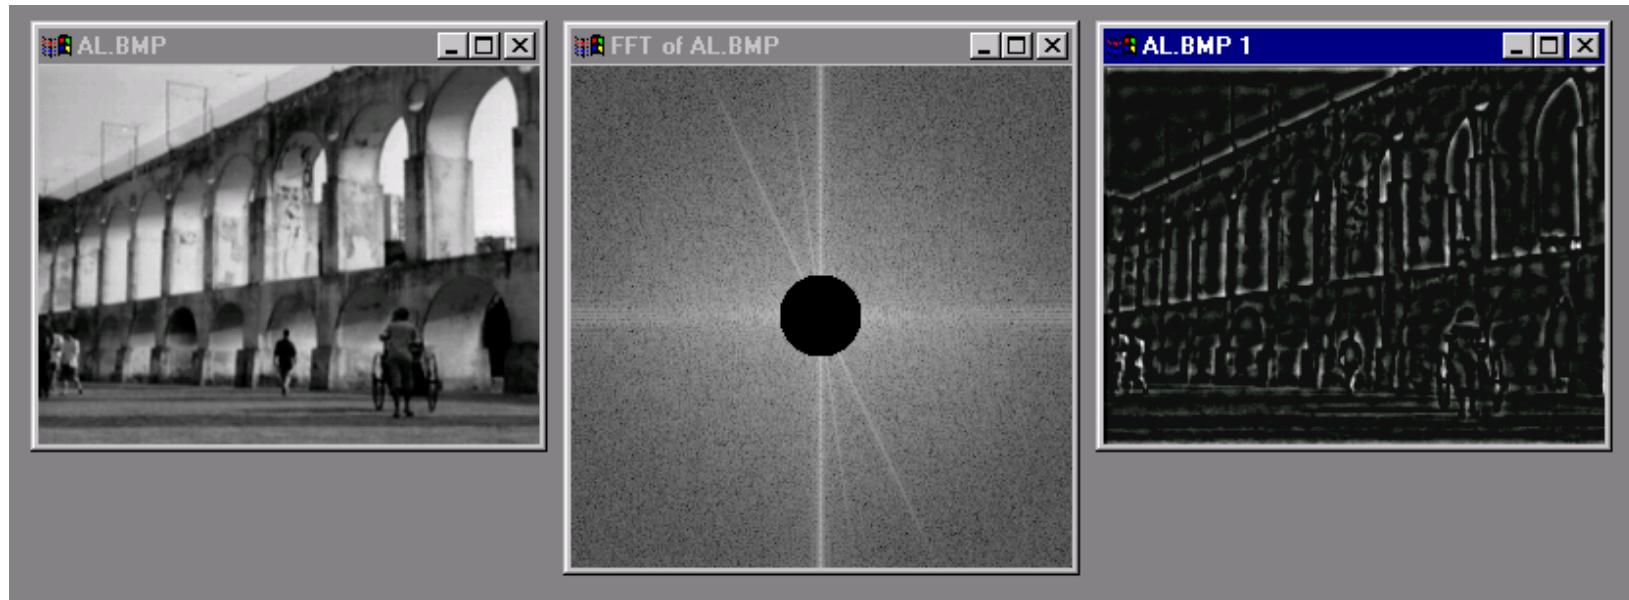
low-pass:



High-pass / band-pass:



Edges in images



Slide credit: A. Efros

Phase and Magnitude

- Curious fact
 - all natural images have about the same magnitude transform
 - hence, phase seems to matter, but magnitude largely doesn't
- Demonstration
 - Take two pictures, swap the phase transforms, compute the inverse - what does the result look like?



Image with cheetah phase
(and zebra magnitude)



Image with zebra phase
(and cheetah magnitude)

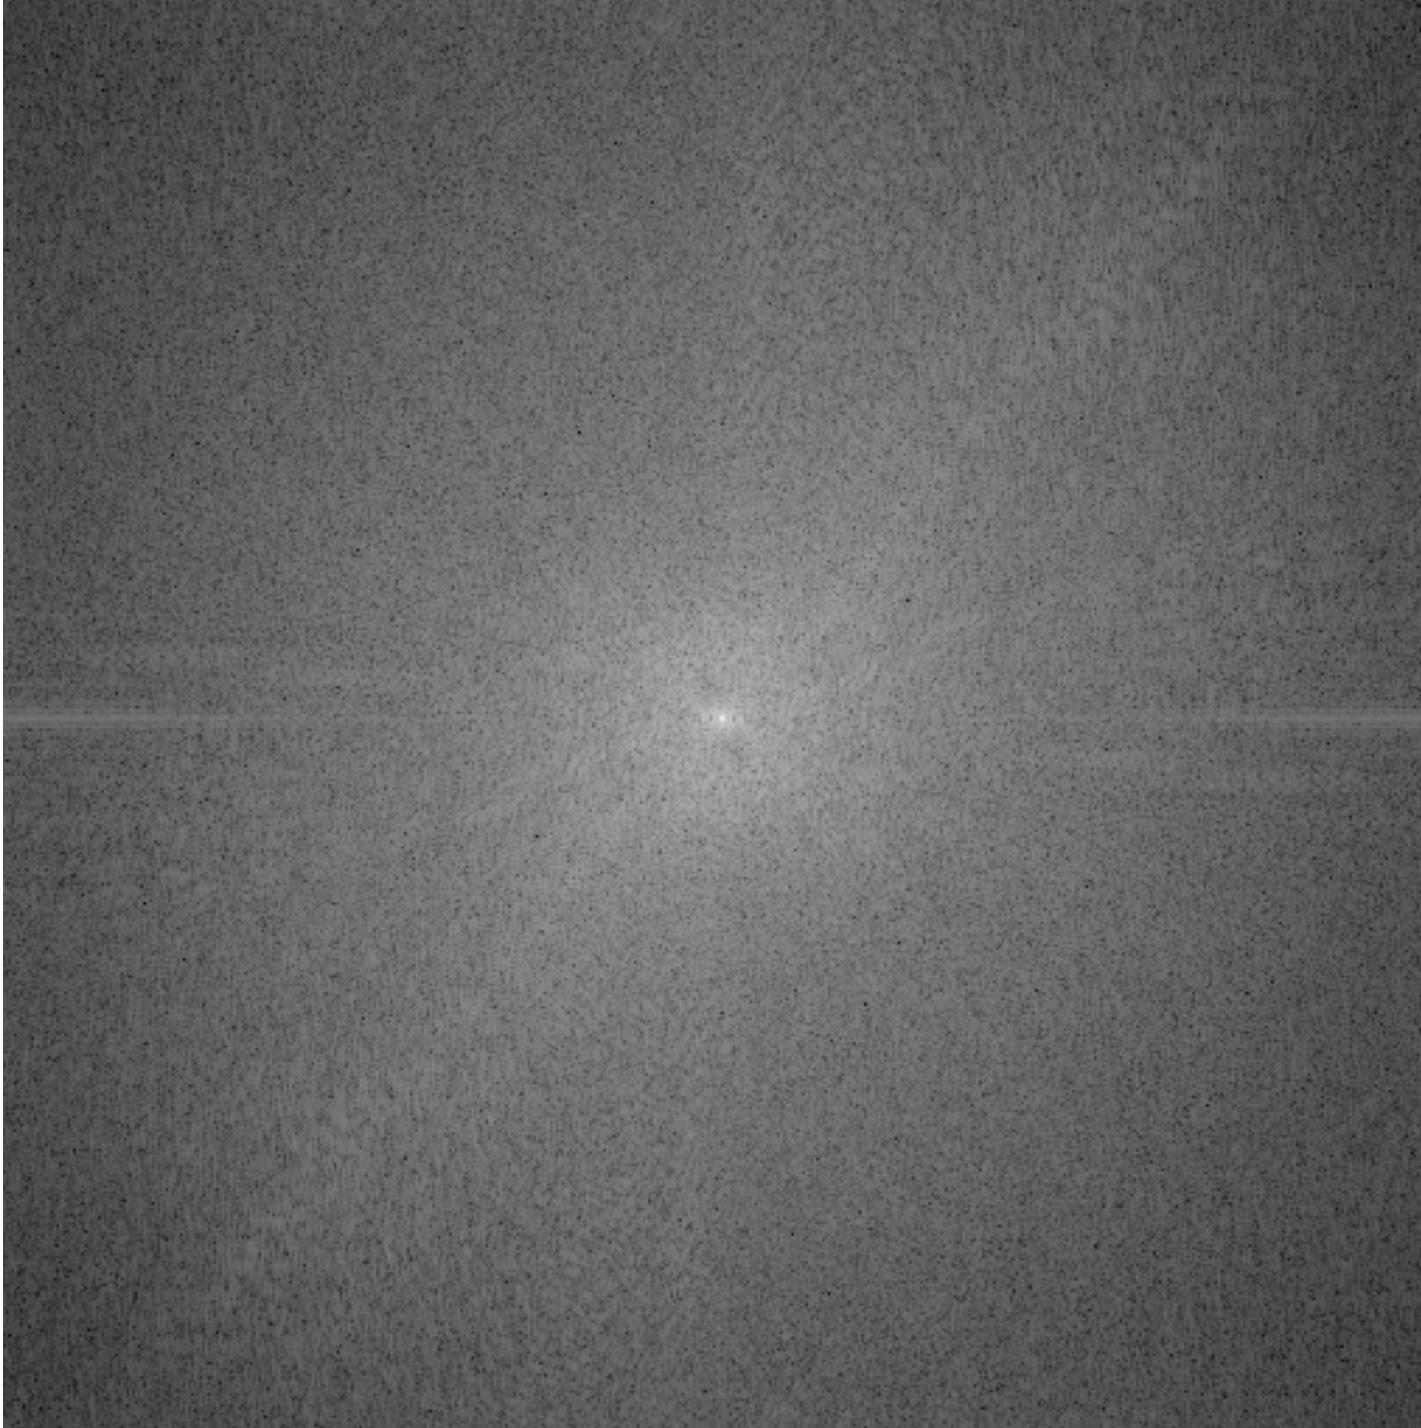


Slide credit: B. Freeman and A. Torralba



Slide credit: B. Freeman and A. Torralba

This is the
magnitude
transform of
the cheetah
picture

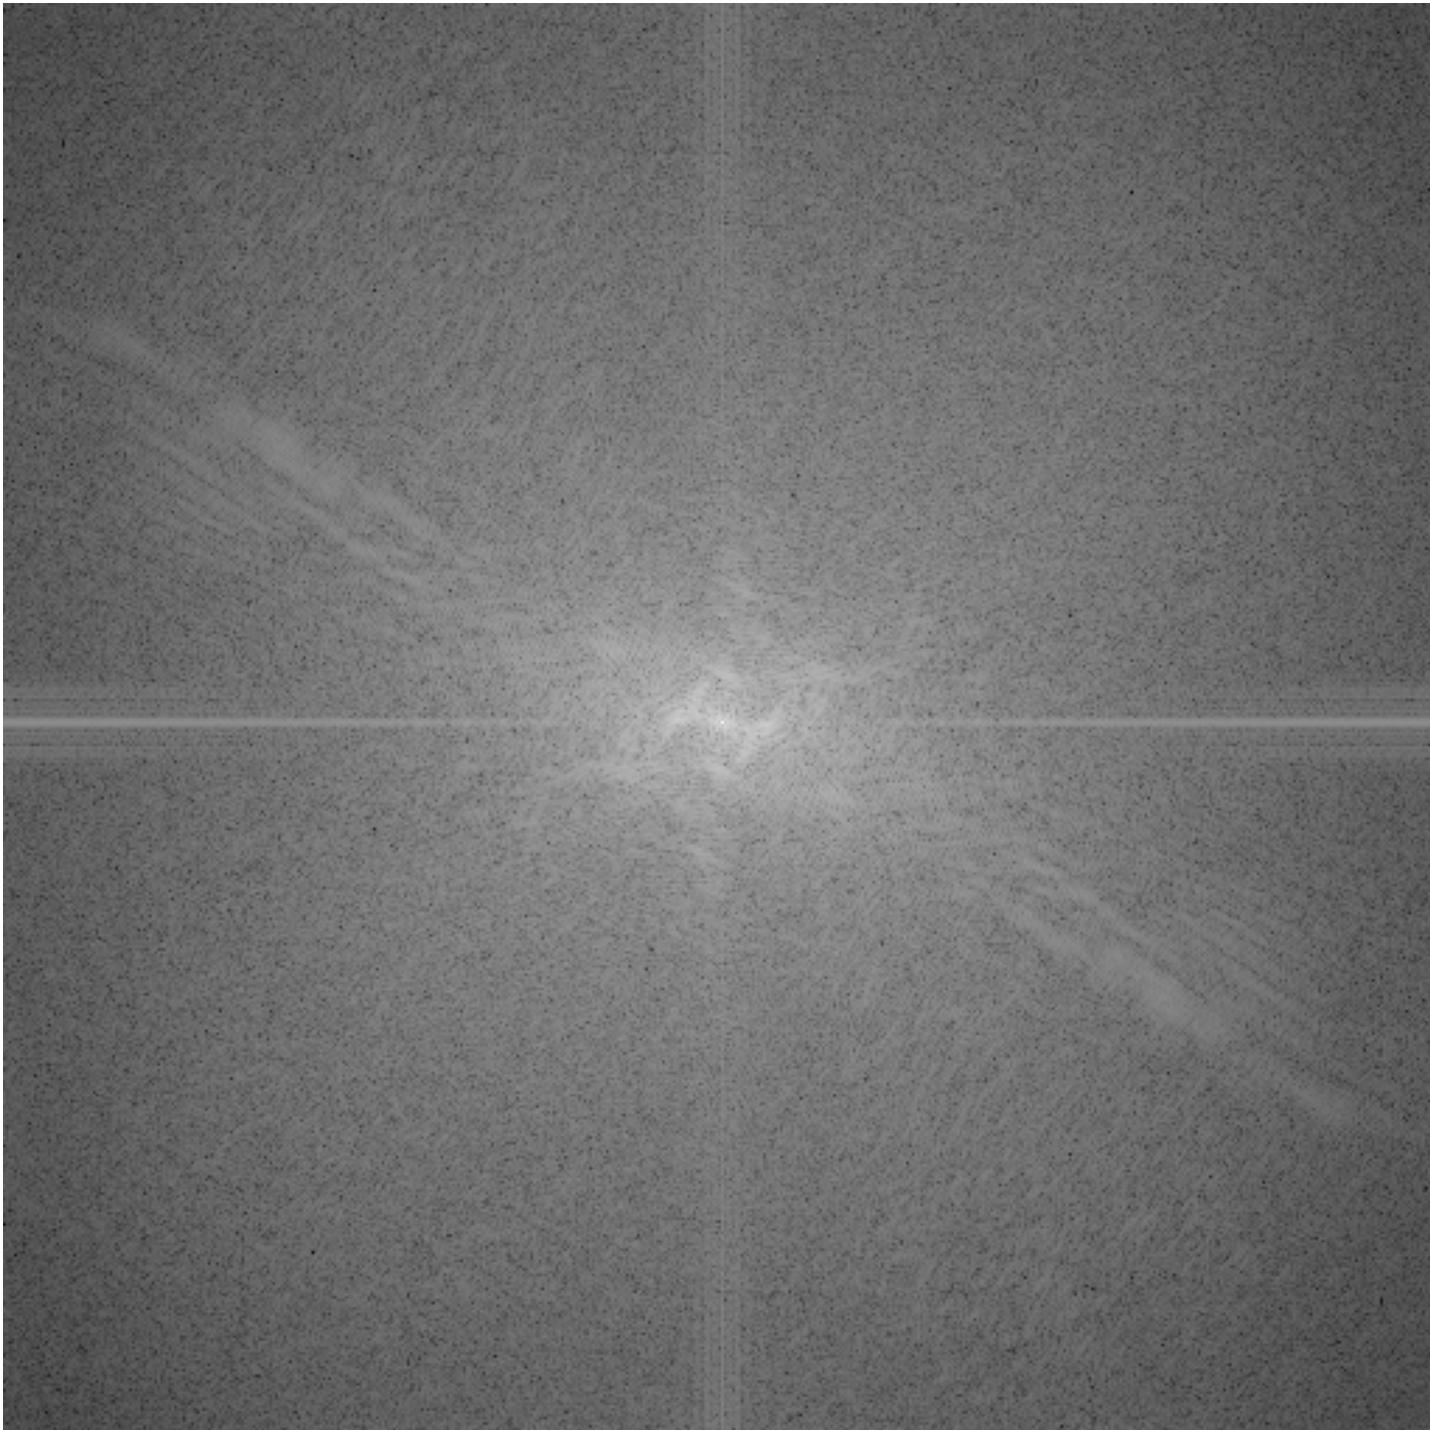


Slide credit: B. Freeman and A. Torralba



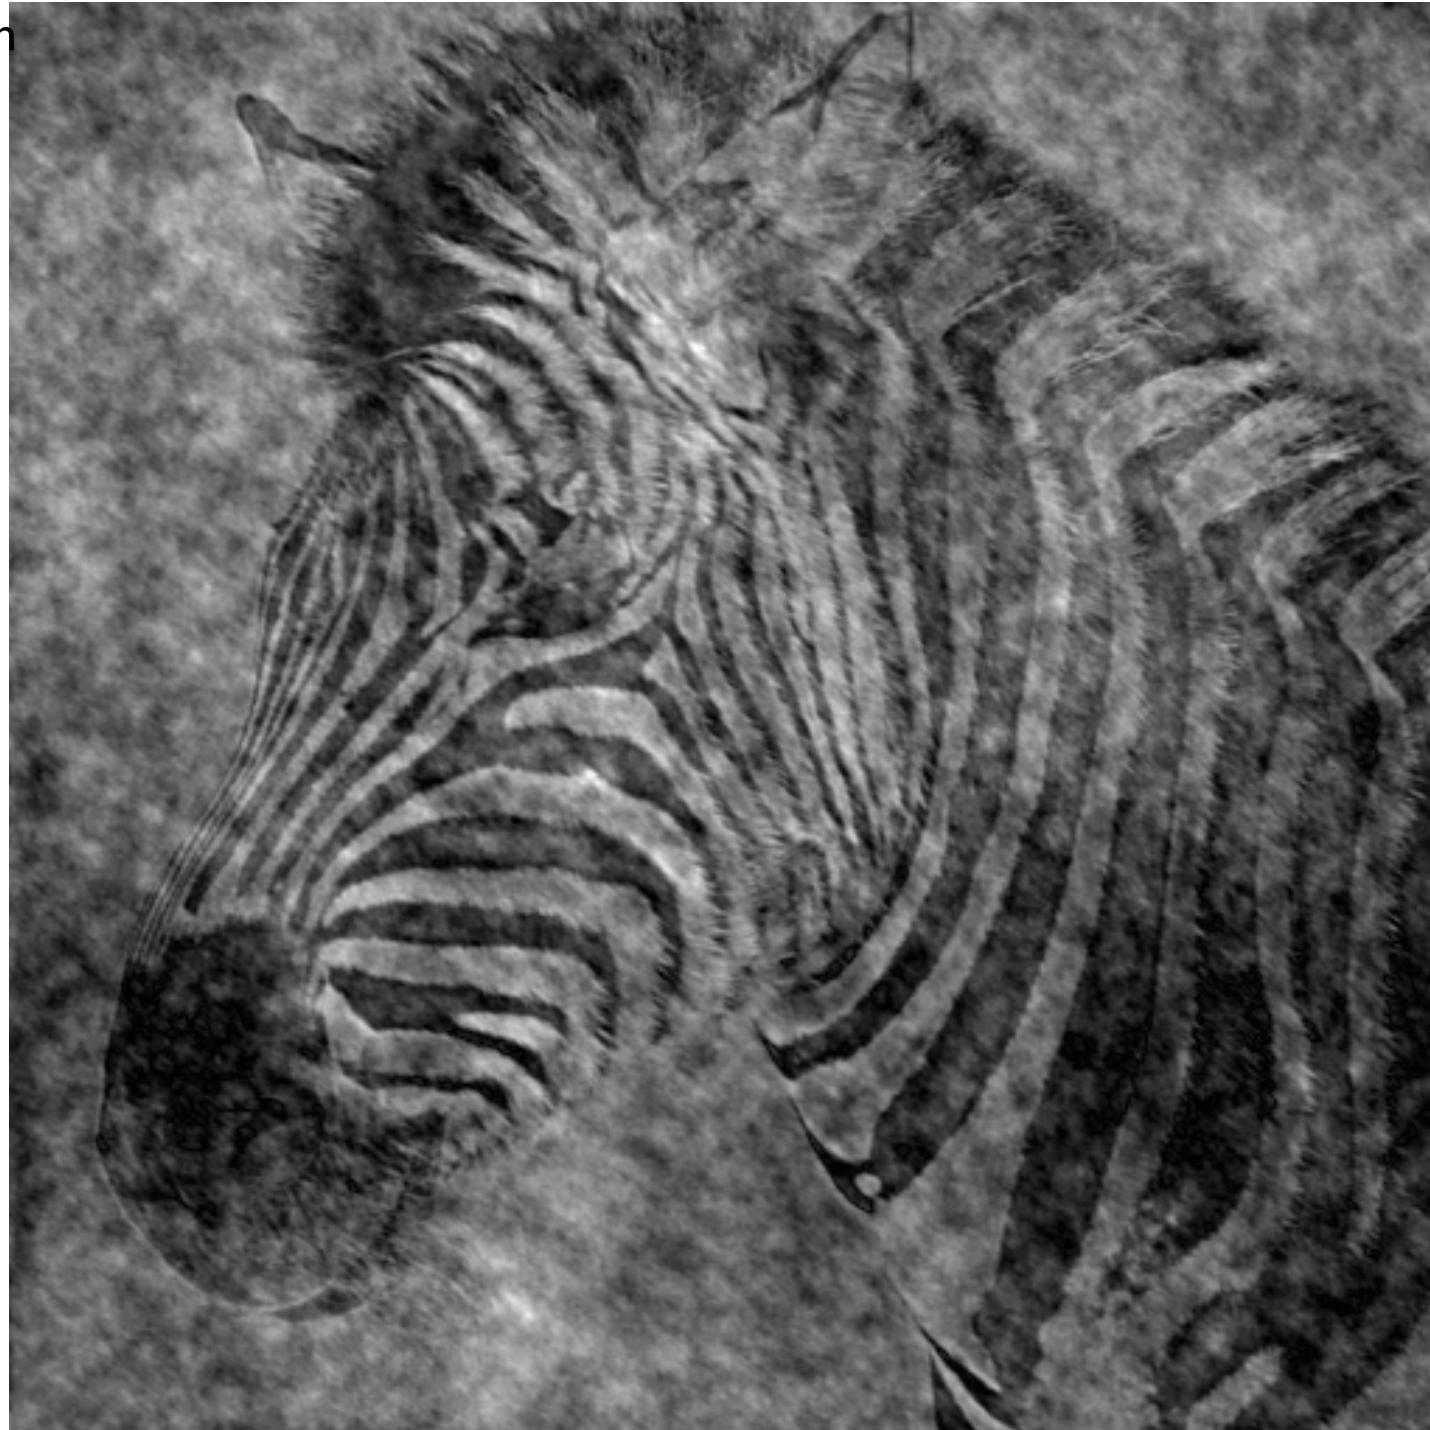
Slide credit: B. Freeman and A. Torralba

This is the
magnitude
transform of
the zebra
picture



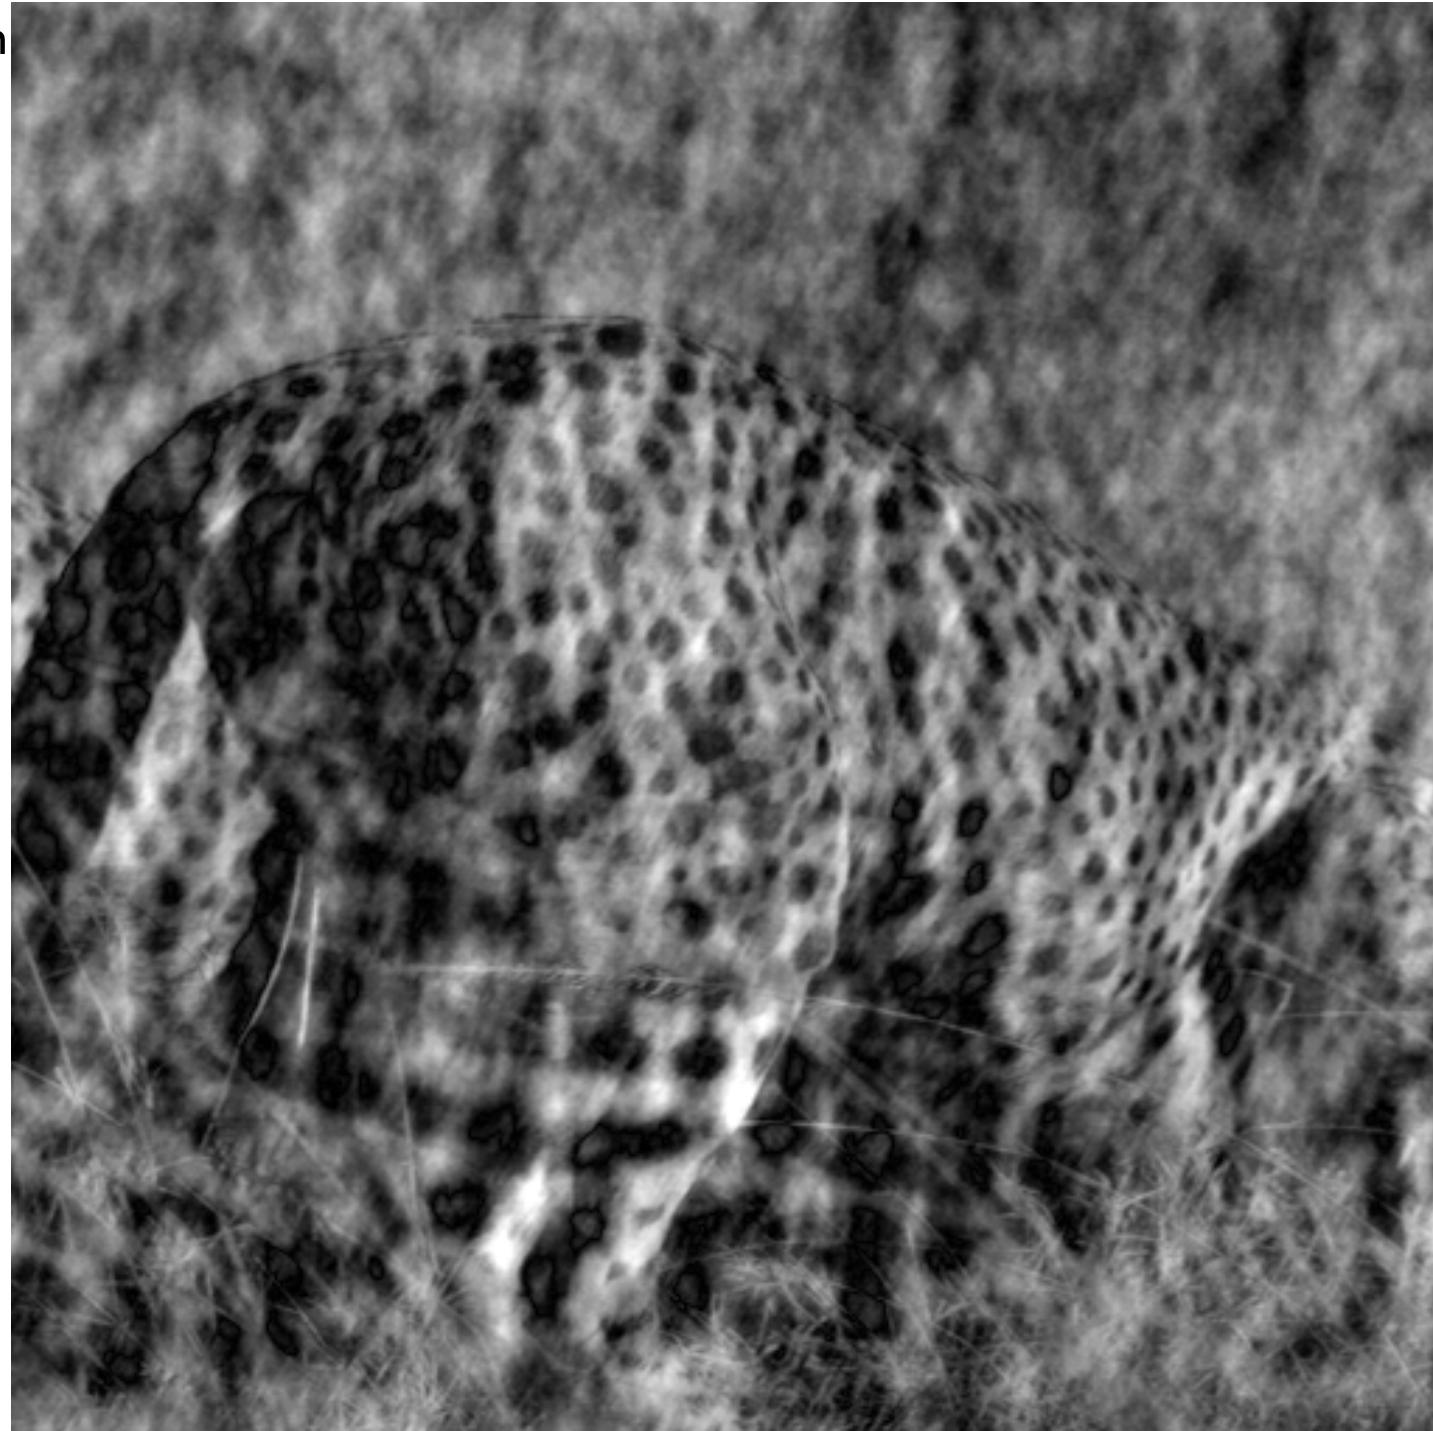
Slide credit: B. Freeman and A. Torralba

Reconstruction
with zebra
phase, cheetah
magnitude



Slide credit: B. Freeman and A. Torralba

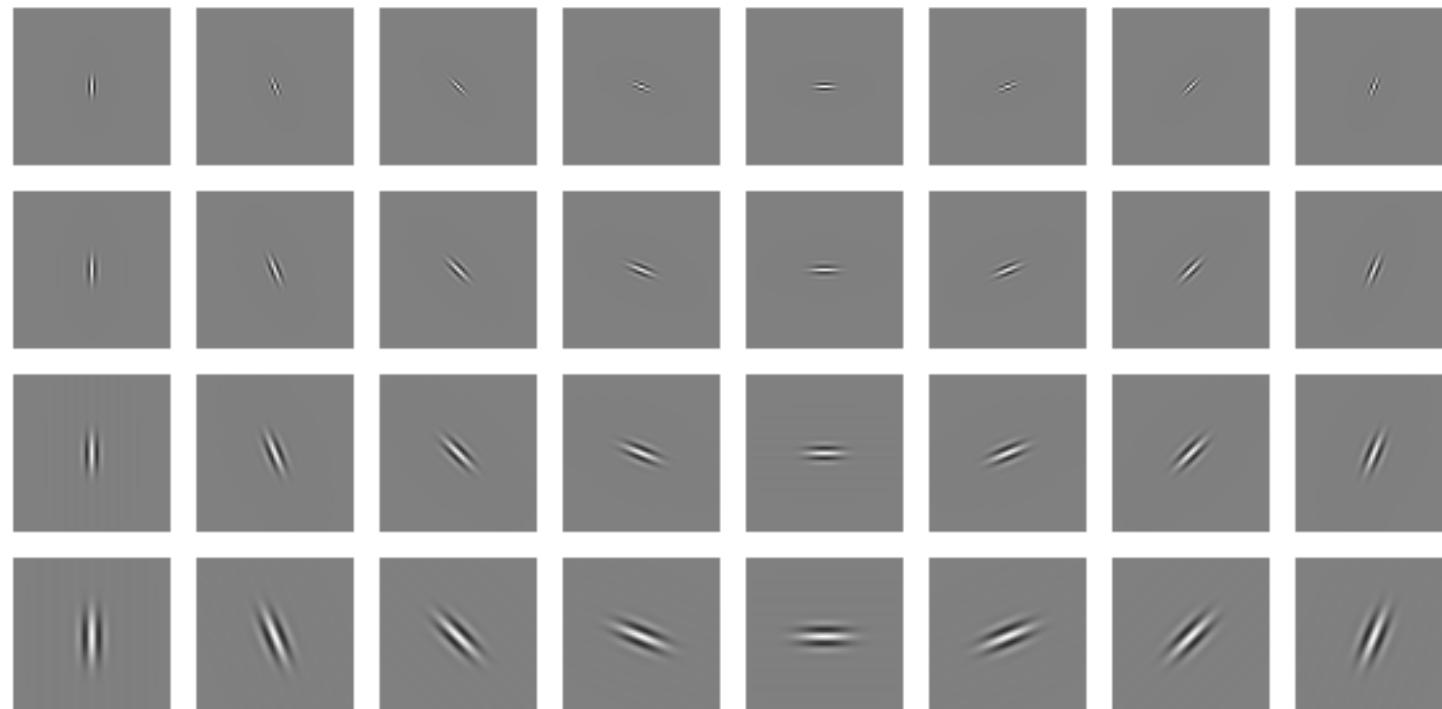
Reconstruction
with cheetah
phase, zebra
magnitude



Slide credit: B. Freeman and A. Torralba

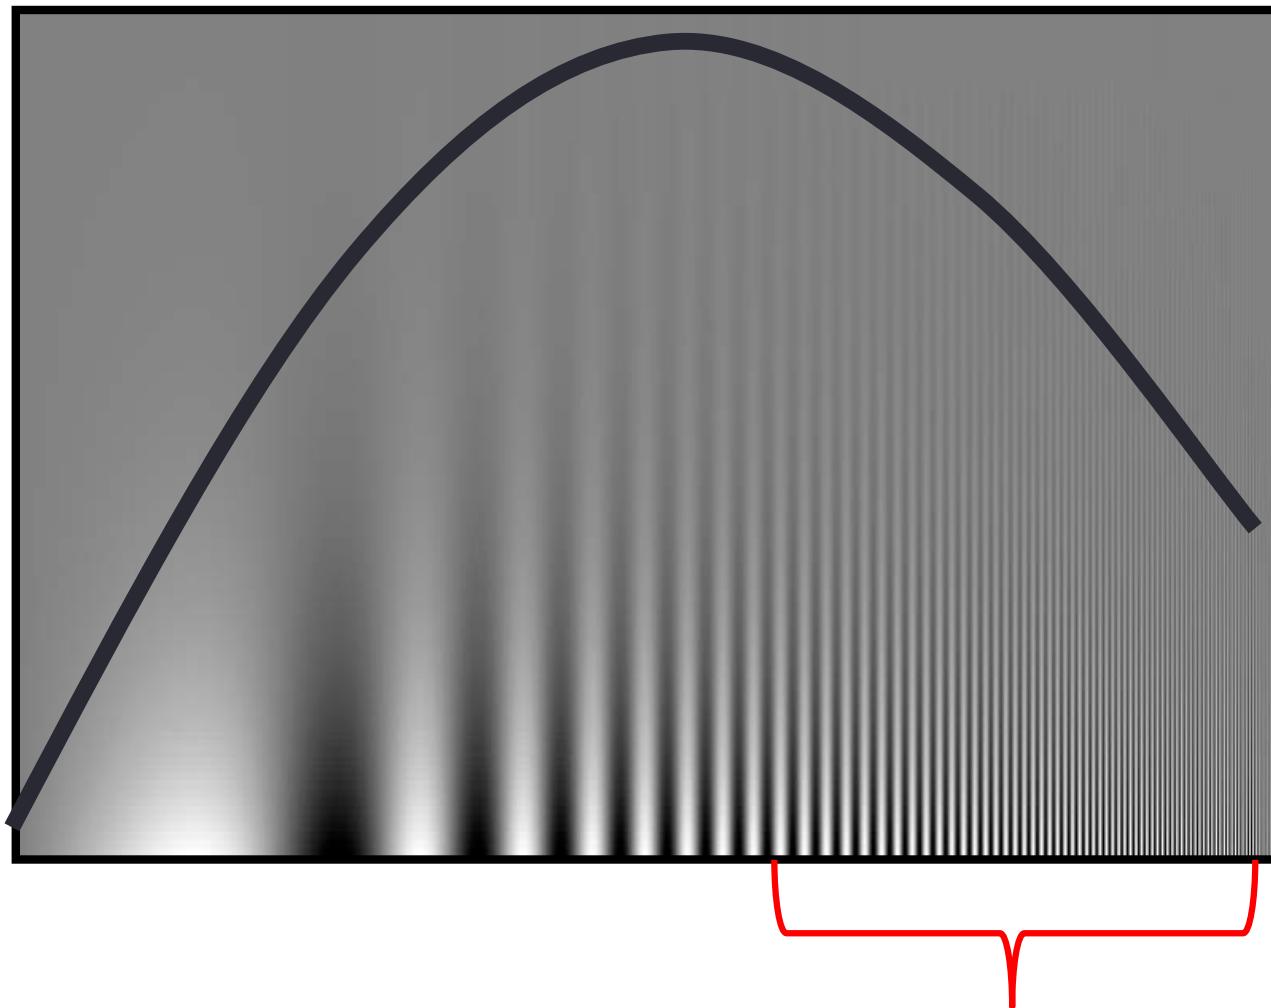
Clues from Human Perception

- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid-high frequencies dominate perception
- When we see an image from far away, we are effectively subsampling it



Early Visual Processing: Multi-scale edge and blob filters

Campbell-Robson contrast sensitivity curve

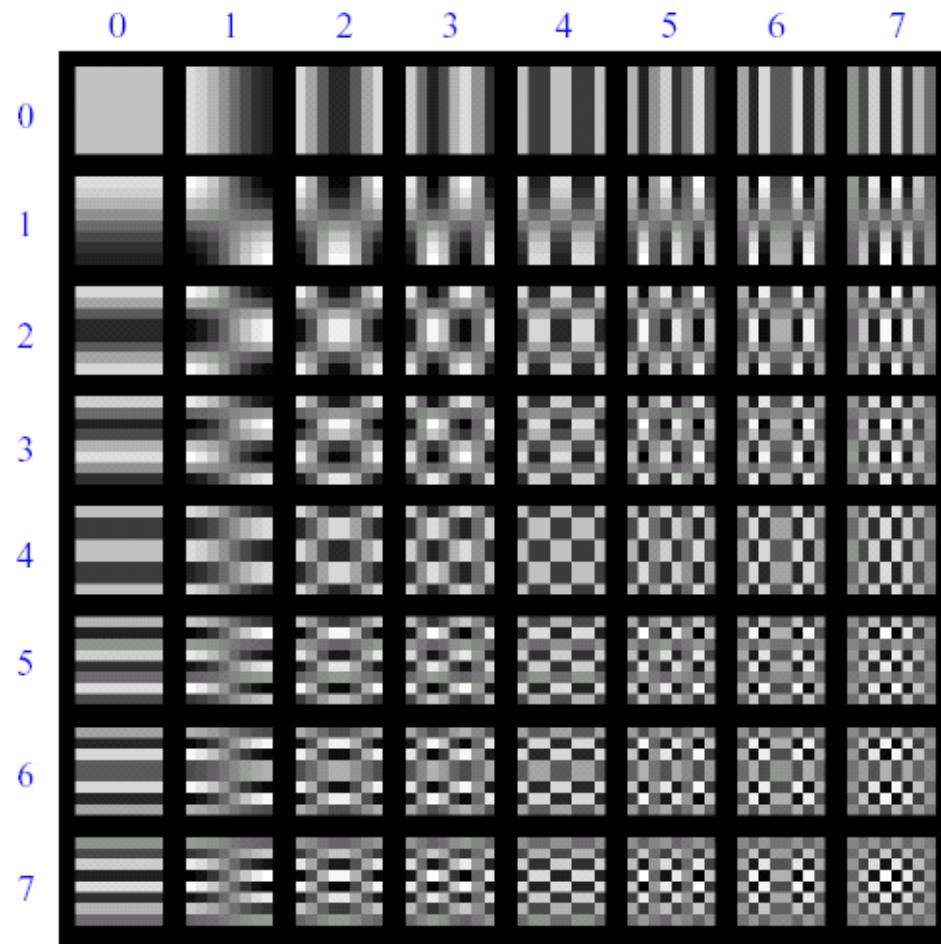
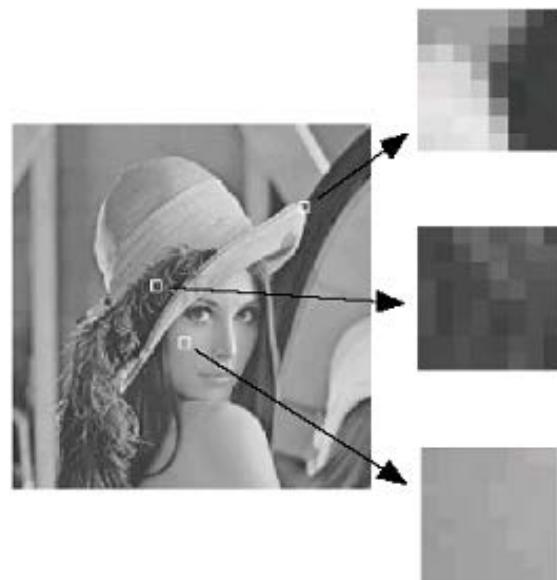


The higher the frequency the less sensitive human visual system is...

Slide credit: J. Hays

Lossy Image Compression (JPEG)

$$X_{k_1, k_2} = \sum_{n_1=0}^{N_1-1} \sum_{n_2=0}^{N_2-1} x_{n_1, n_2} \cos \left[\frac{\pi}{N_1} \left(n_1 + \frac{1}{2} \right) k_1 \right] \cos \left[\frac{\pi}{N_2} \left(n_2 + \frac{1}{2} \right) k_2 \right].$$



Block-based Discrete Cosine Transform (DCT) on 8x8 Slide credit: A. Bobick

Using DCT in JPEG

- The first coefficient $B(0,0)$ is the DC component, the average intensity
- The top-left coeffs represent low frequencies, the bottom right – high frequencies

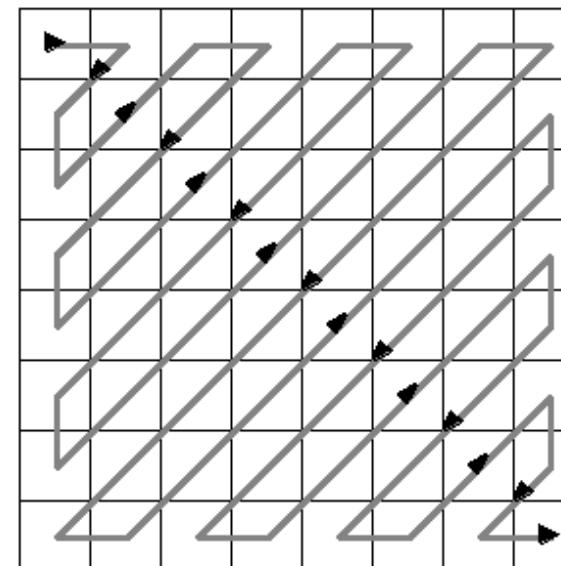
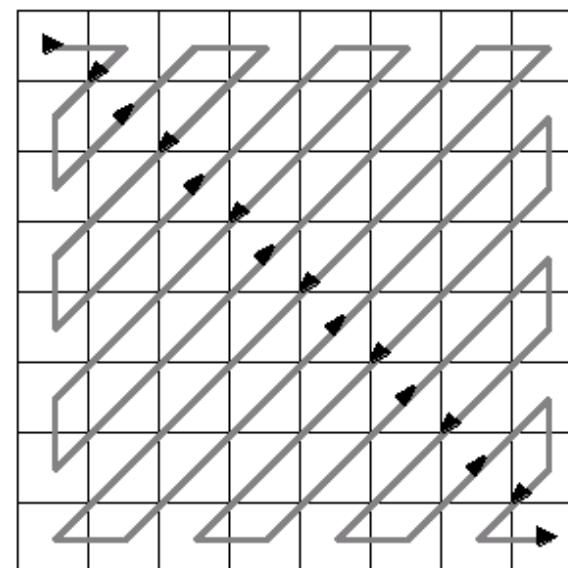


Image compression using DCT

- DCT enables image compression by concentrating most image information in the low frequencies
- Loose unimportant image info (high frequencies) by cutting $B(u,v)$ at bottom right
- The decoder computes the inverse DCT – IDCT



JPEG compression comparison



89k



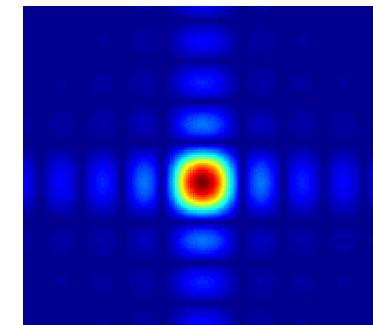
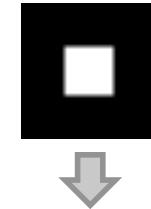
12k

Slide credit: A. Bobick

Things to Remember

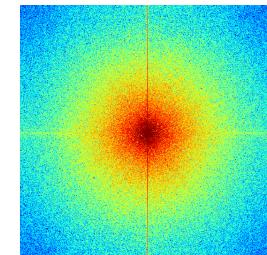
- Sometimes it makes sense to think of images and filtering in the frequency domain

- Fourier analysis



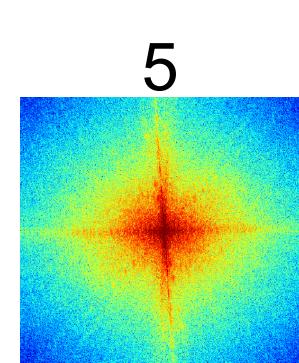
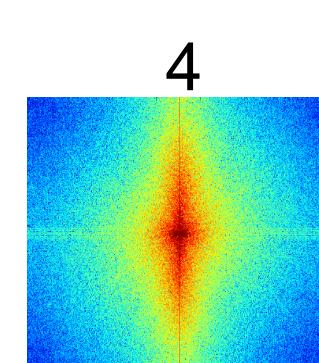
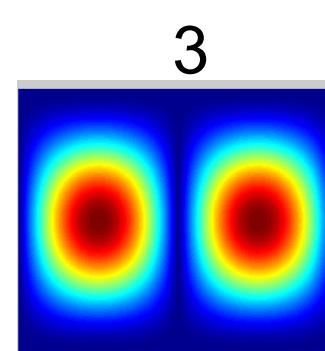
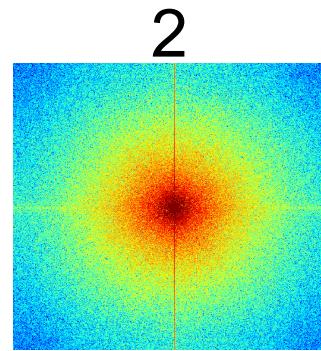
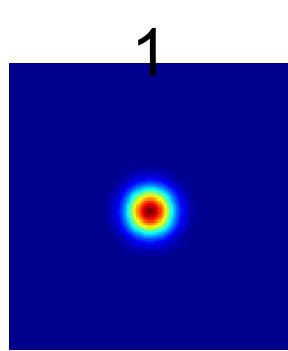
- Can be faster to filter using FFT for large images ($N \log N$ vs. N^2 for auto-correlation)

- Images are mostly smooth
 - Basis for compression



Practice question

1. Match the spatial domain image to the Fourier magnitude image



A

